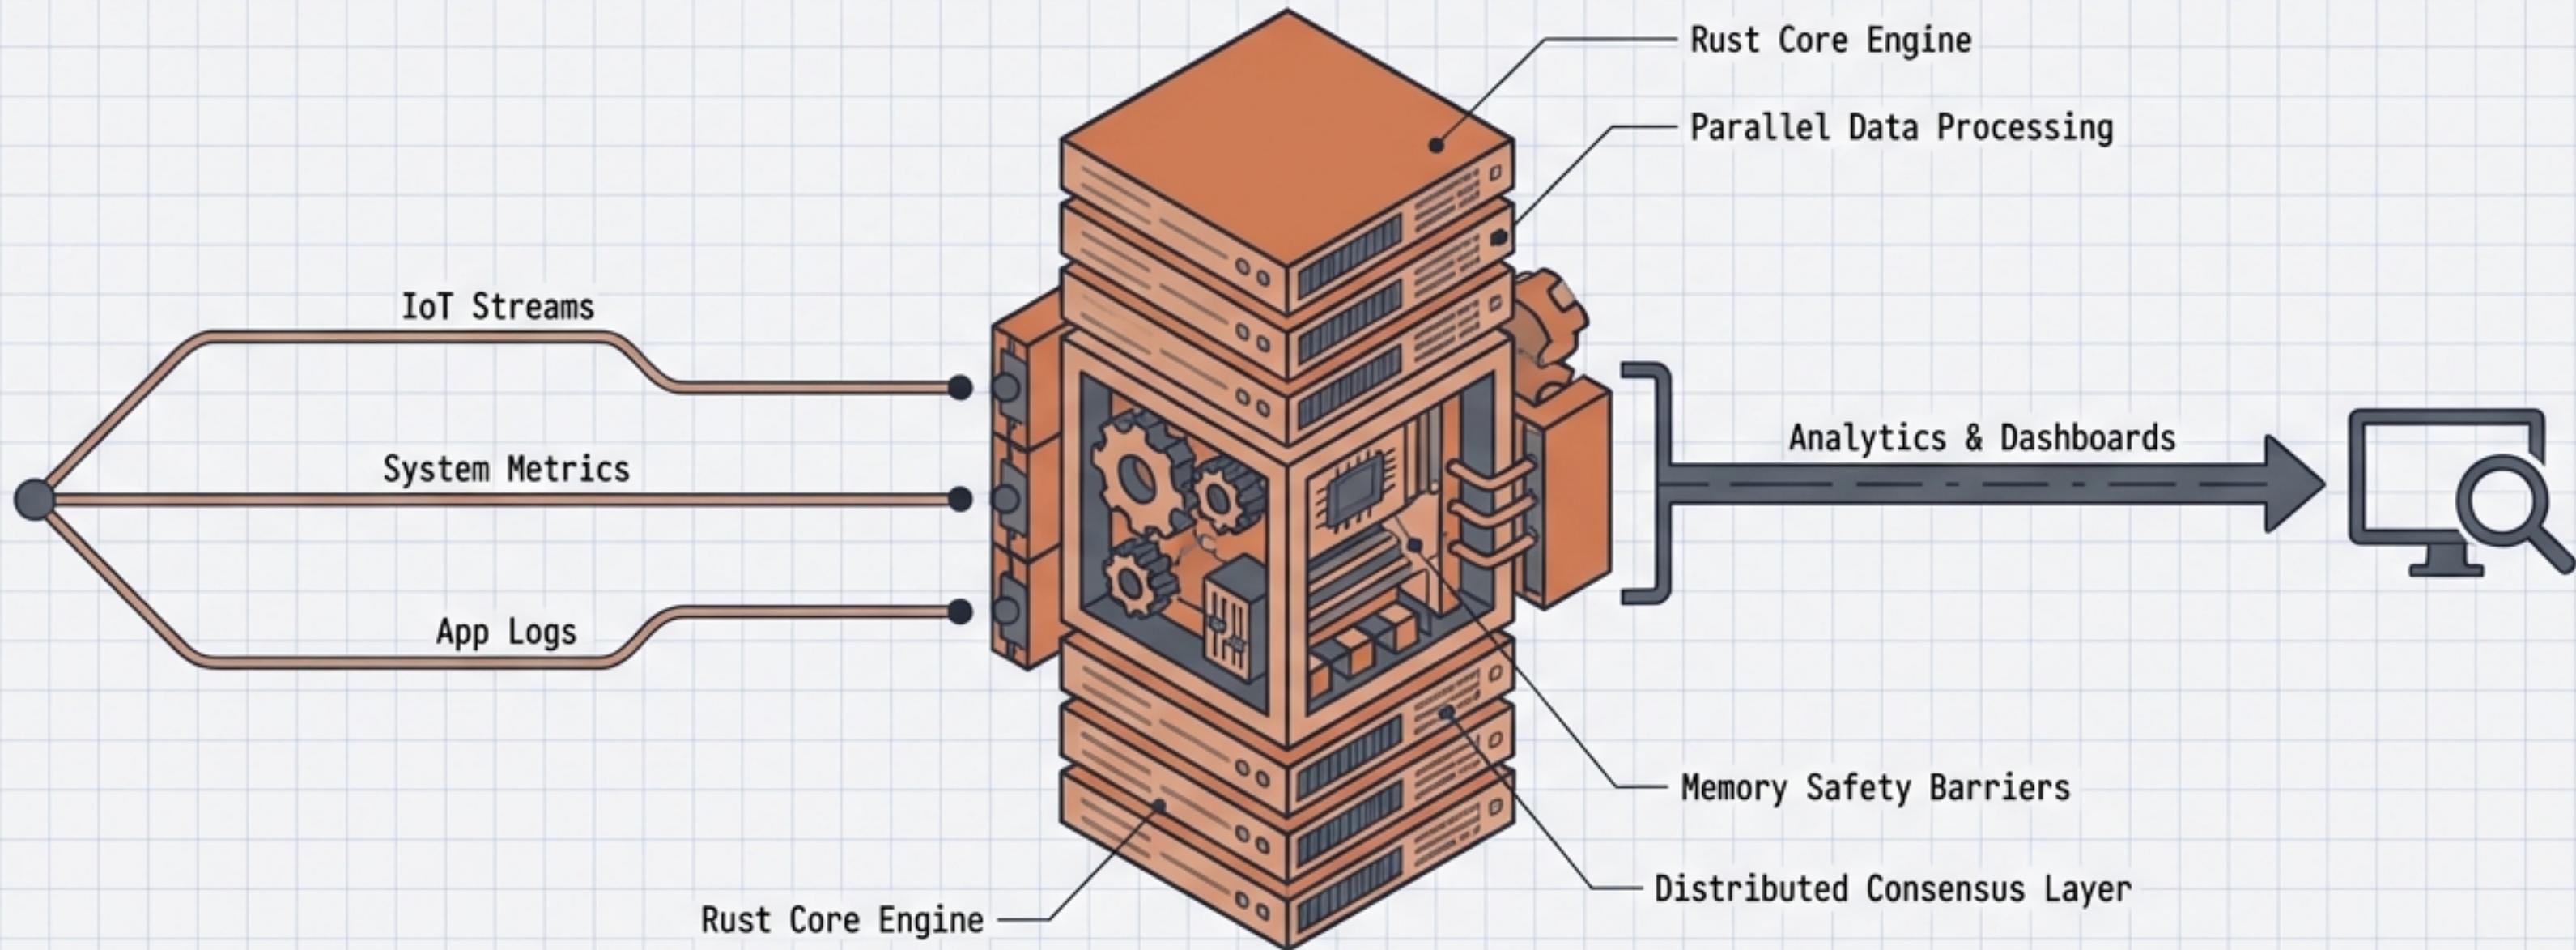


RusTs: Precision Engineering for Time Series Data

A high-performance, memory-safe, distributed database built for the velocity of IoT and real-time analytics.



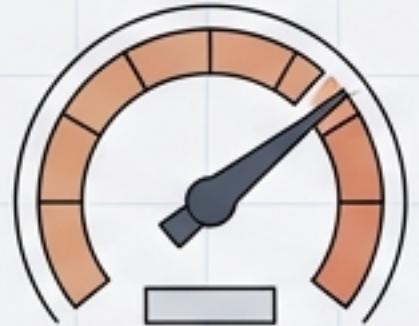
Implementation: Native Rust 1.75+

Protocol: InfluxDB Line Compatible

Storage: Columnar + Gorilla/Delta Compression

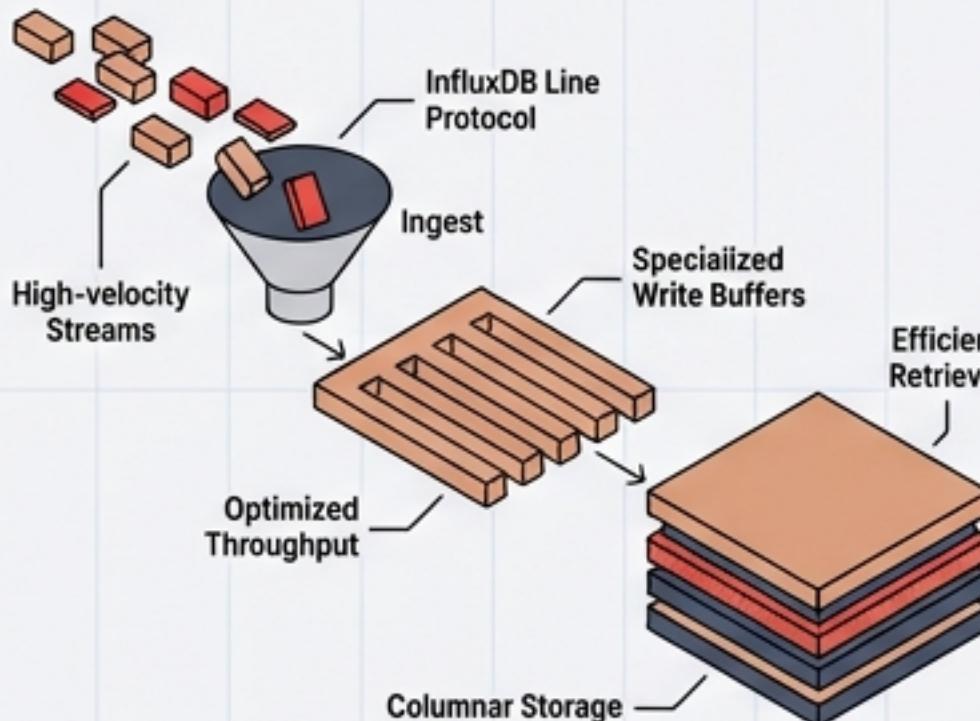
License: Dual Apache 2.0 / MIT

Engineered for Speed, Efficiency, and Reliability



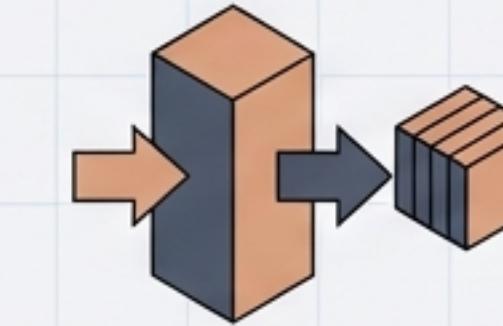
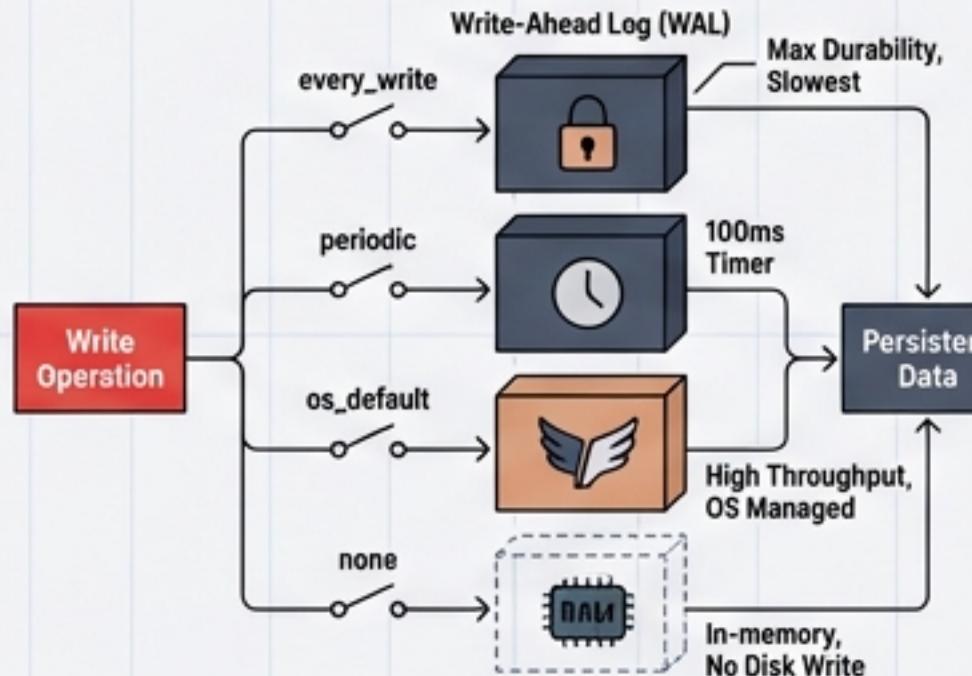
Performance & Compatibility

- Optimized for high-velocity time series workloads using columnar storage.
- Drop-in migration: Supports InfluxDB line protocol for immediate ecosystem integration.
- High-throughput ingestion via specialized write buffers.



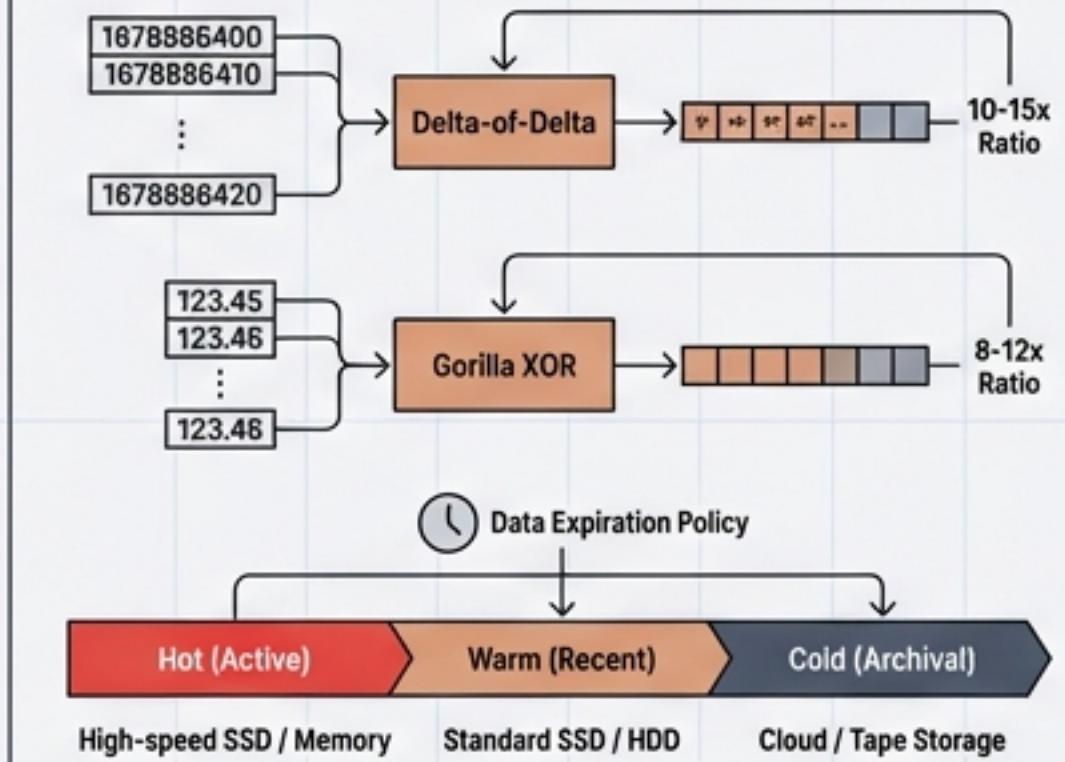
Flexible Durability

- Configurable Write-Ahead Log (WAL) modes allow tuning the consistency-performance trade-off.
- Modes:
 - `every_write` (Sync, max durability)
 - `periodic` (Default 100ms sync)
 - `os_default` (High throughput)
 - `none` (In-memory/Testing)

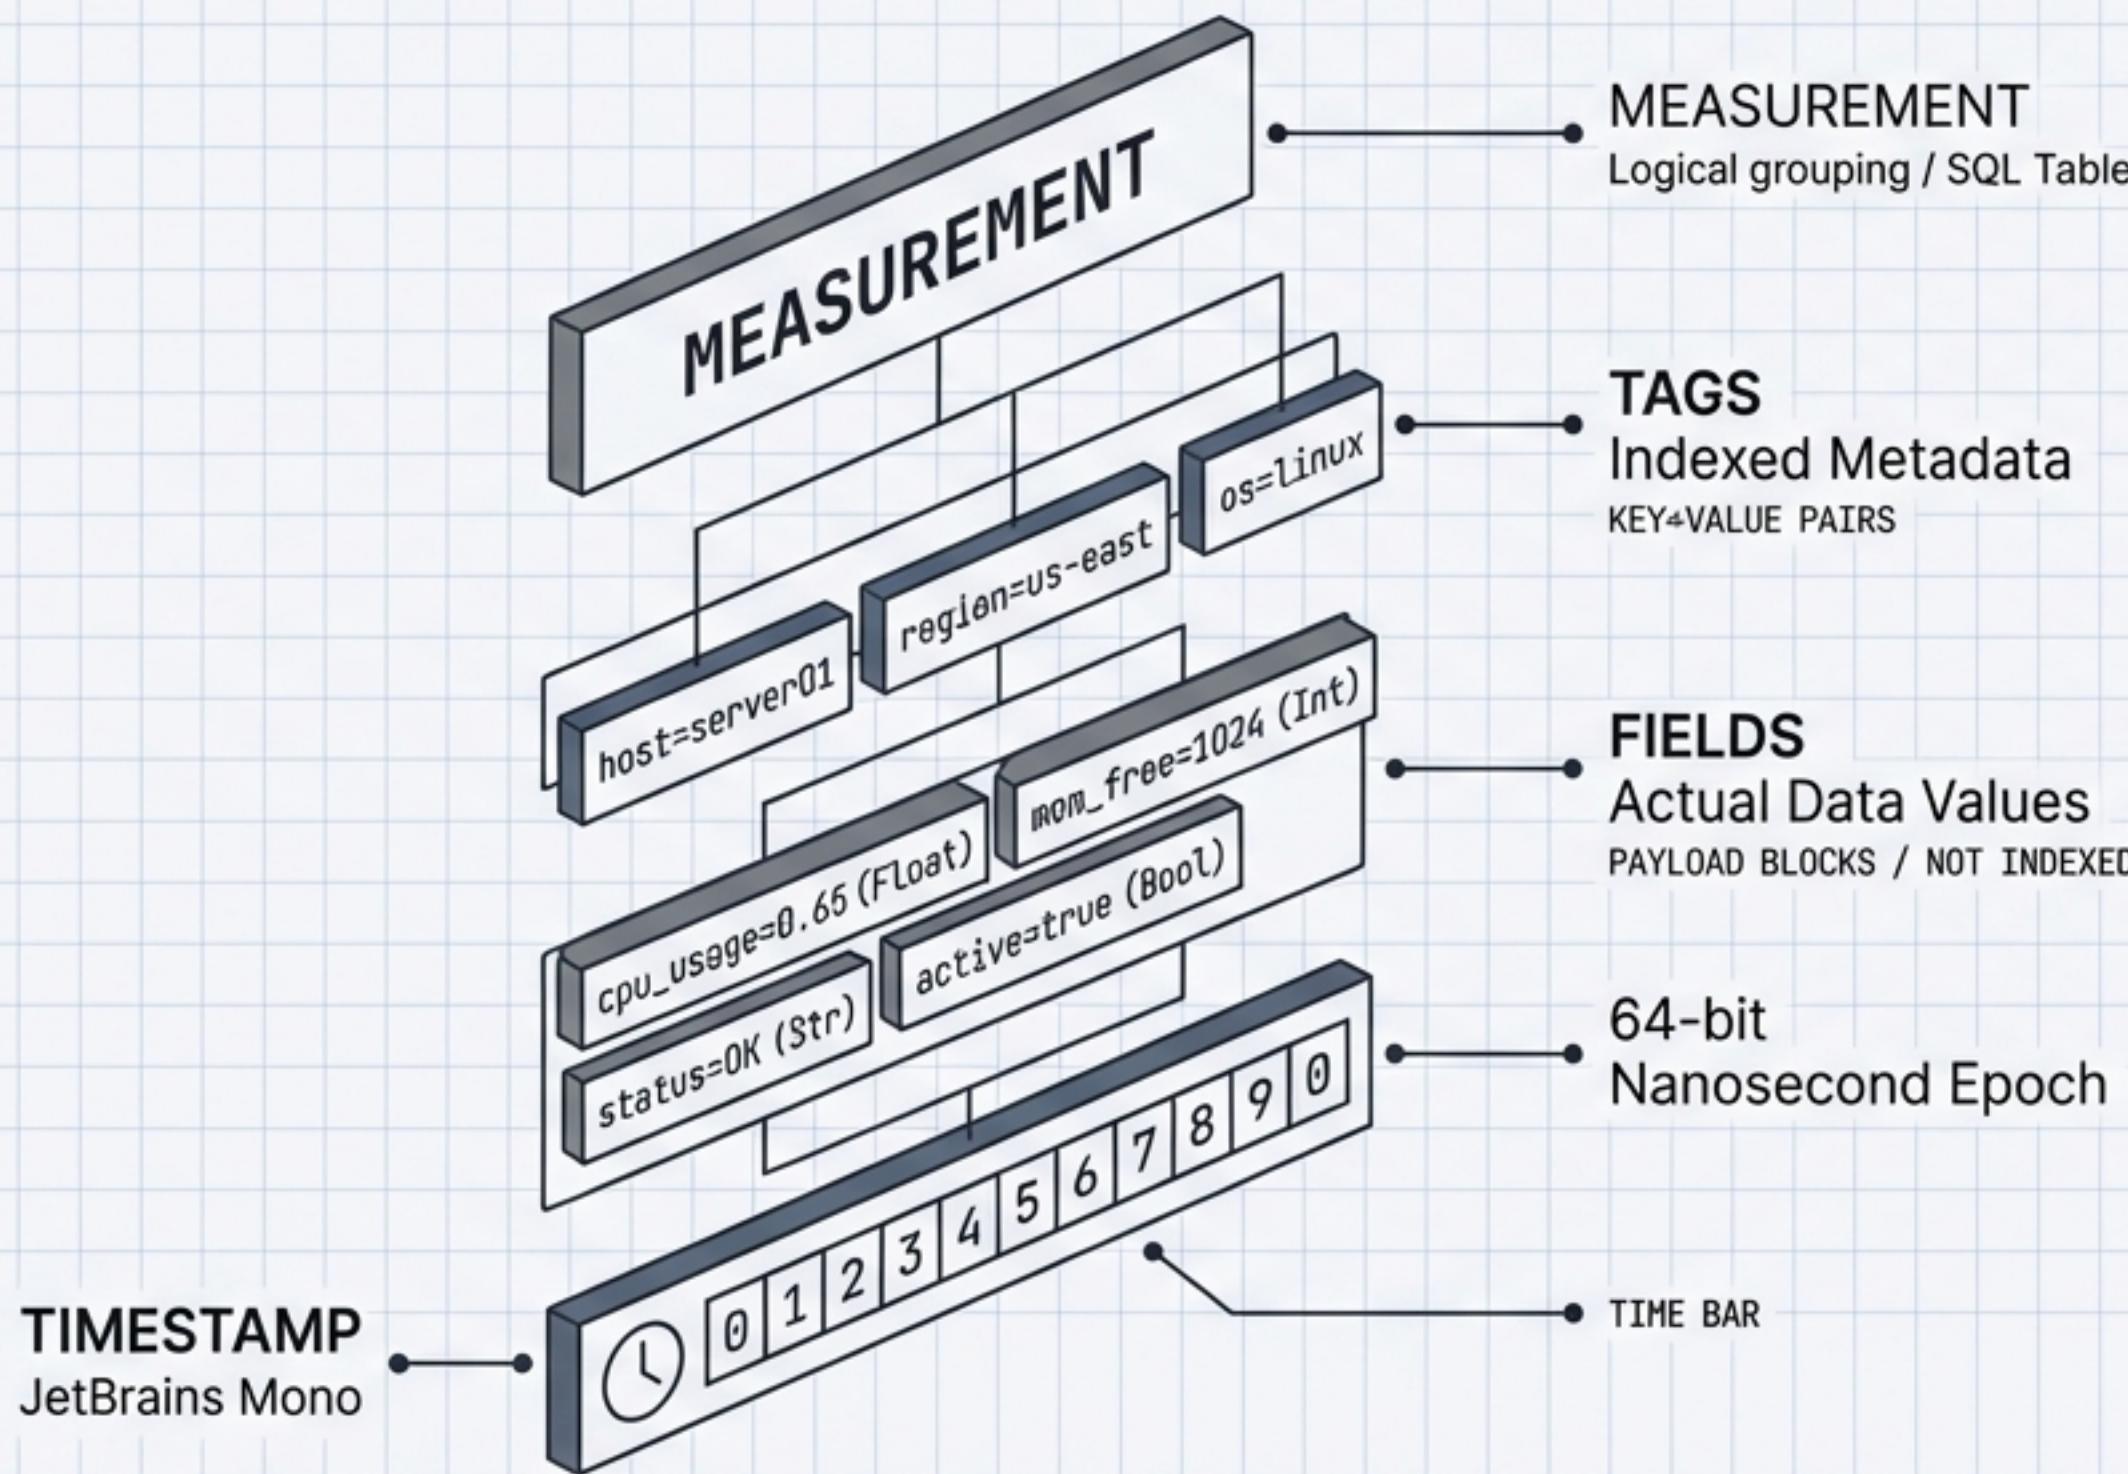


Advanced Optimization

- 10-15x compression ratios on timestamps using Delta-of-Delta.
- 8-12x compression on floats using Gorilla XOR.
- Automatic data expiration and storage tiering (Hot/Warm/Cold).

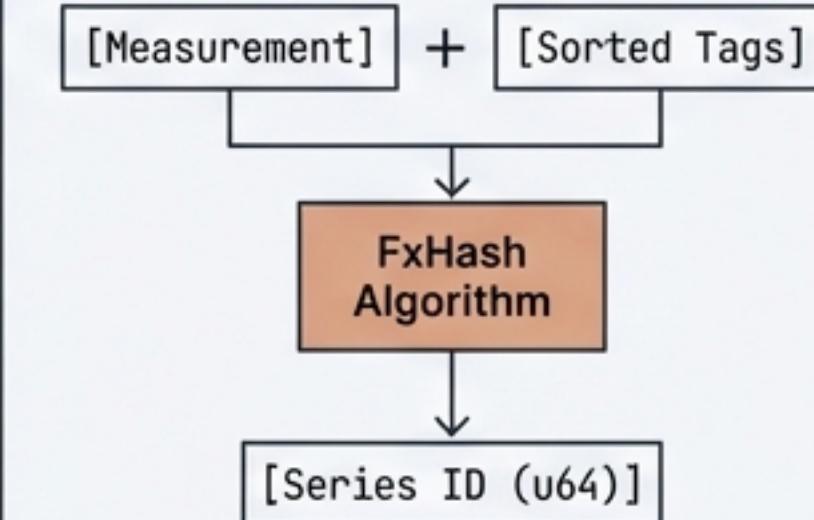


The Atomic Structure of Time Series Data



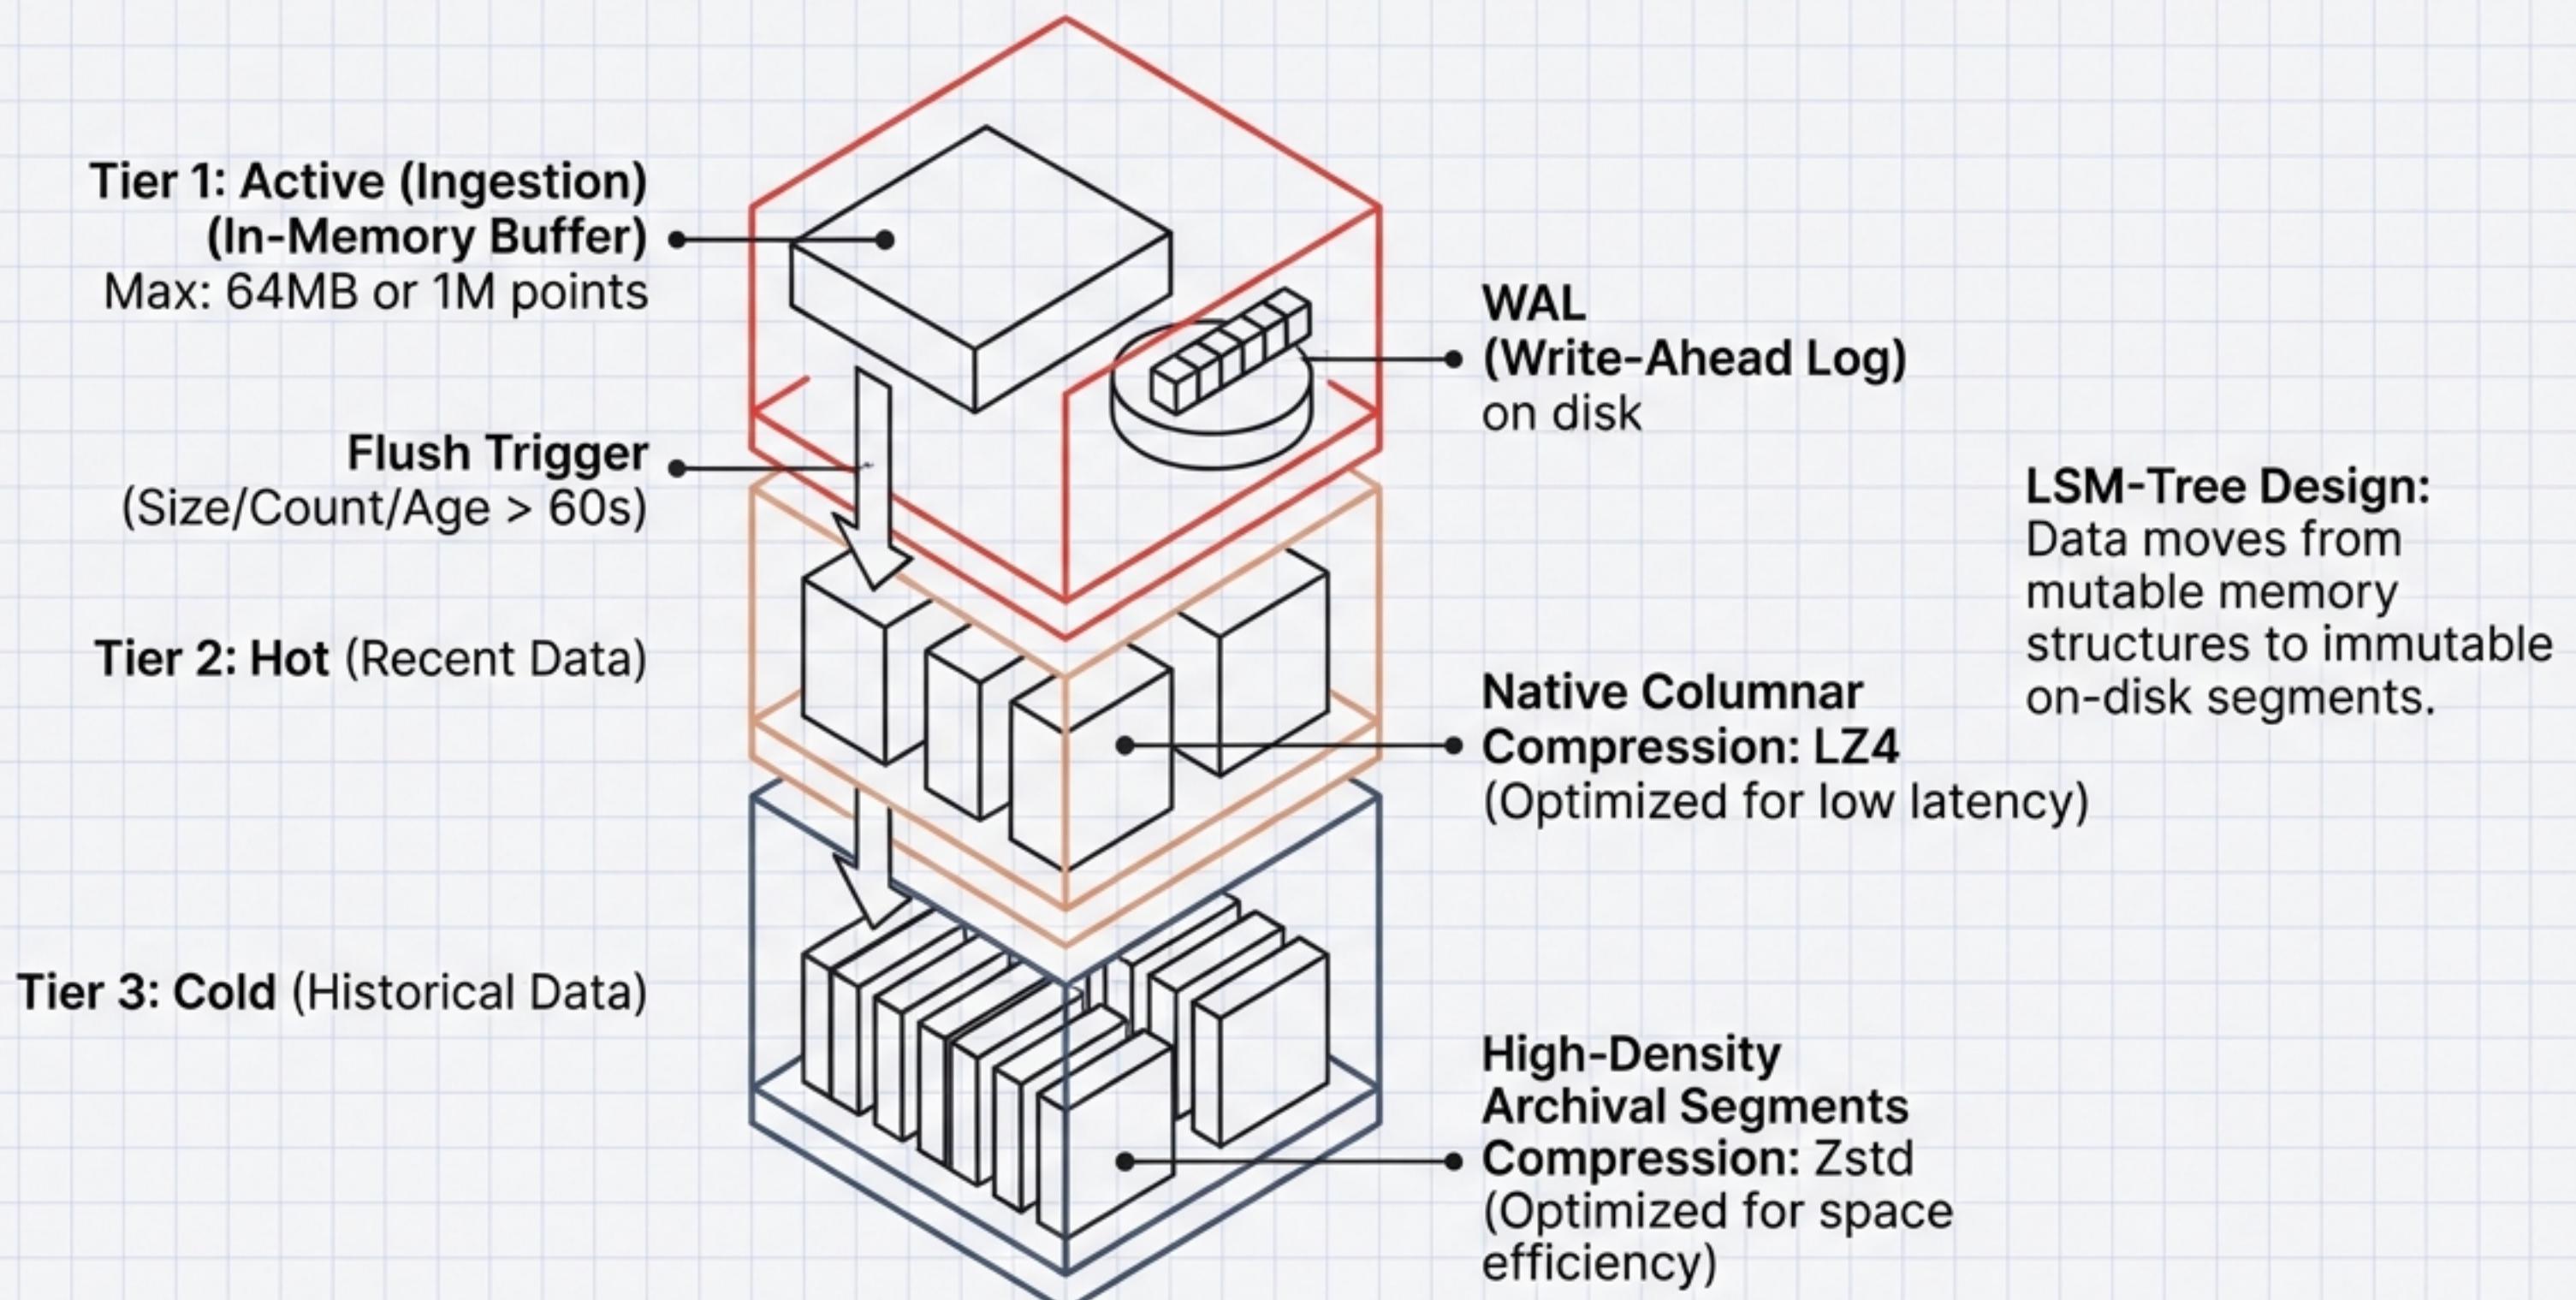
The Series ID (u64)

A unique identifier representing a specific time series.

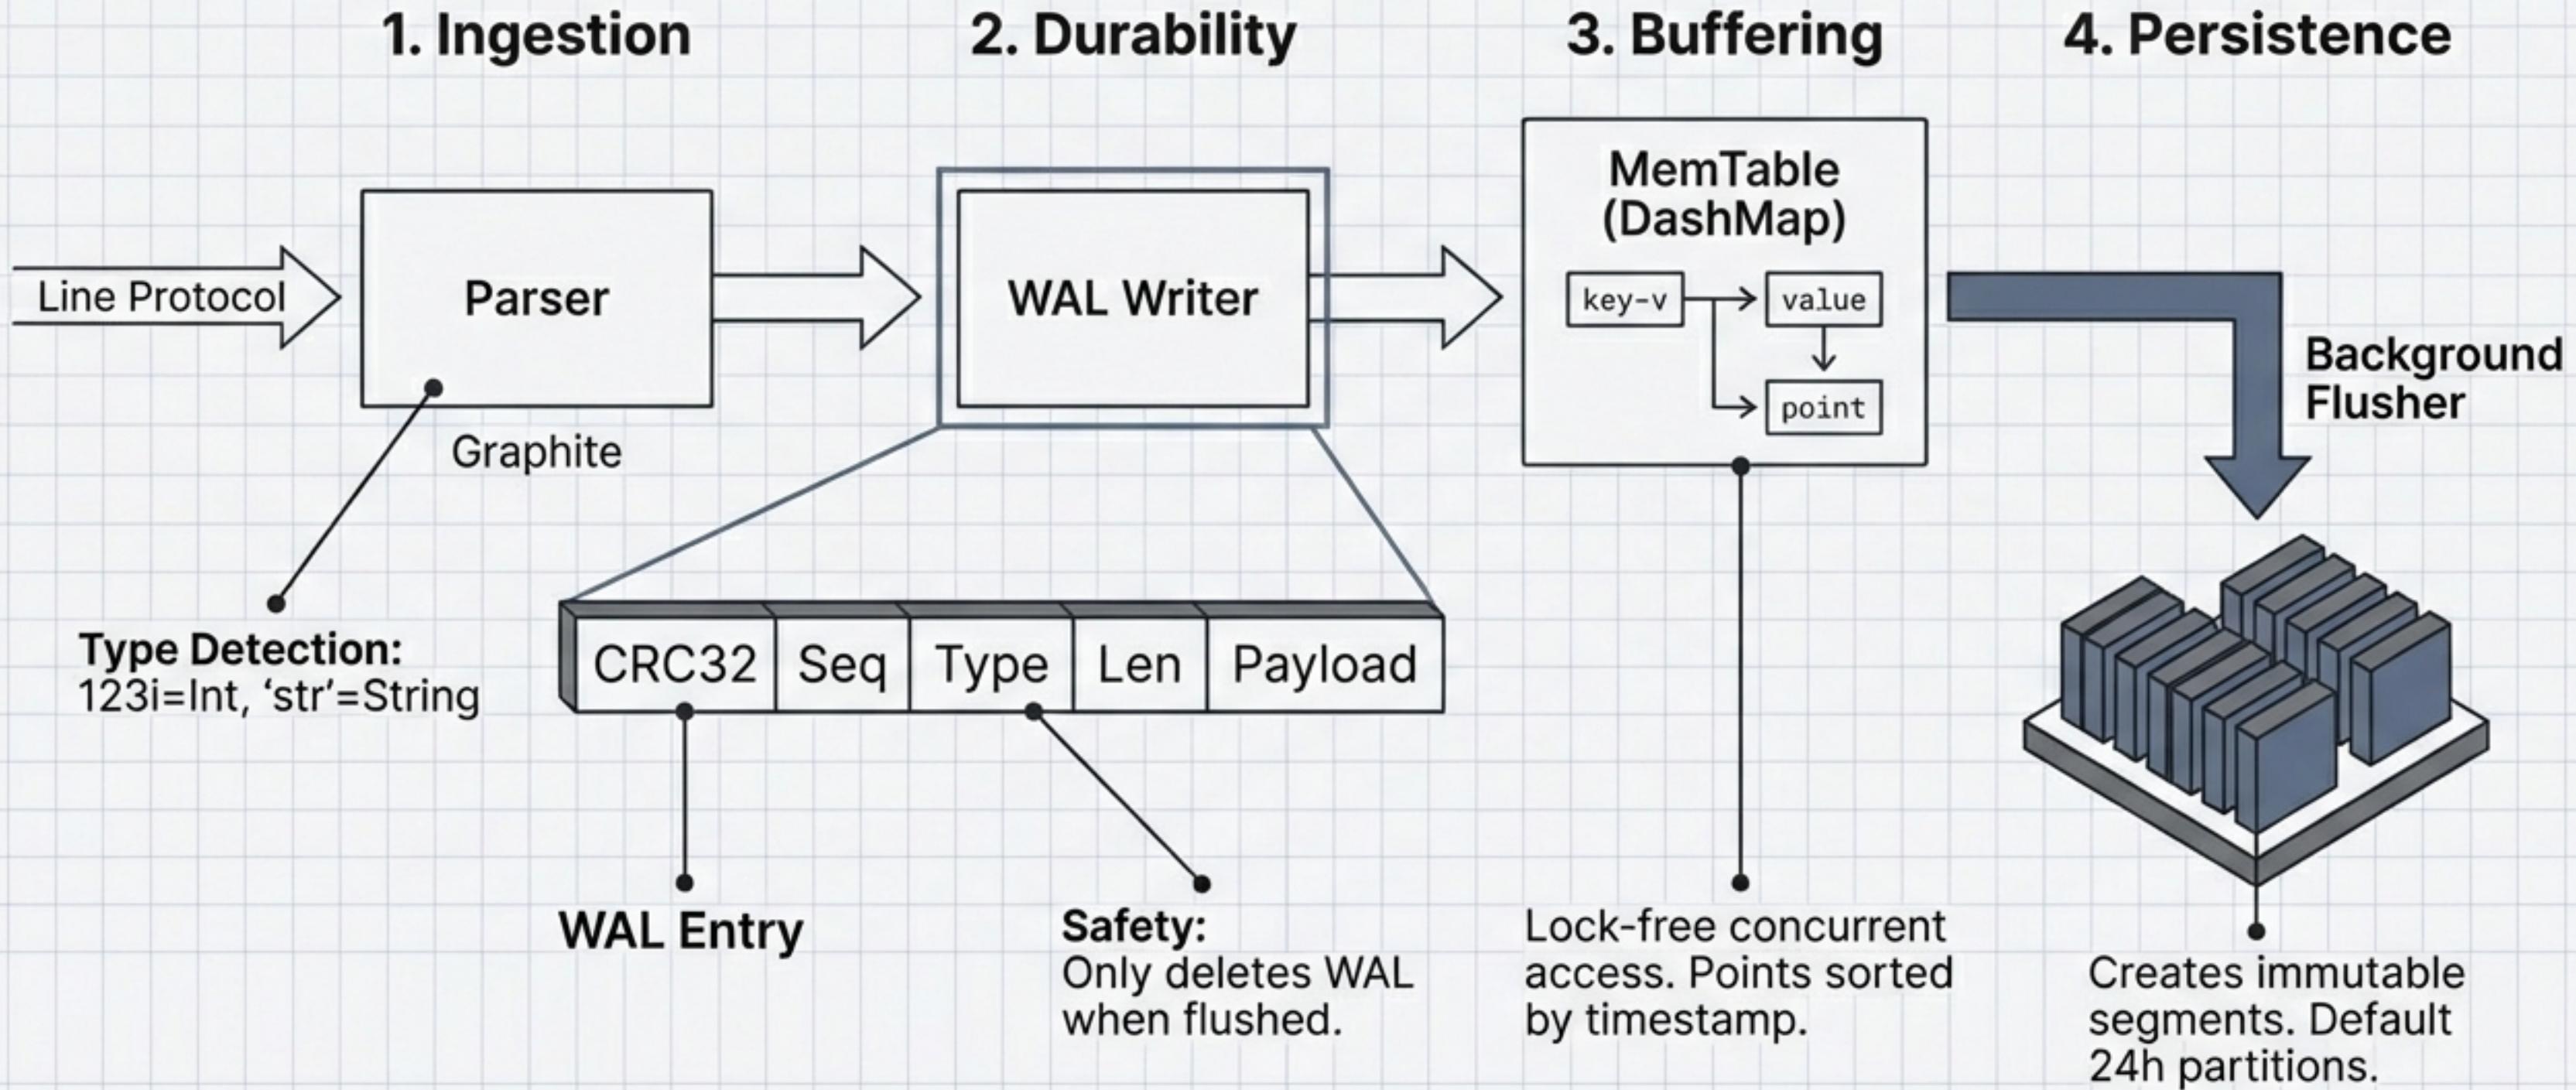


Constraint: Tags are always sorted alphabetically before hashing to ensure deterministic IDs.

A Three-Tiered Storage Architecture

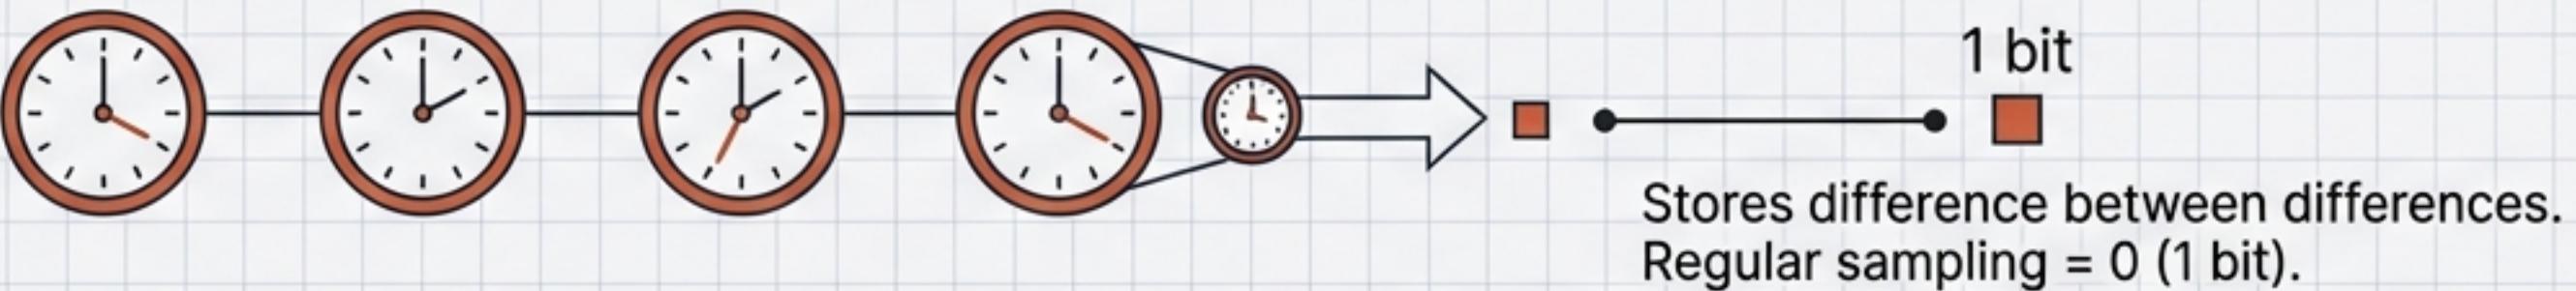


The Write Path and Durability Guarantees



Algorithmic Efficiency: The Compression Layer

Timestamps
Delta-of-Delta
10-15x Ratio

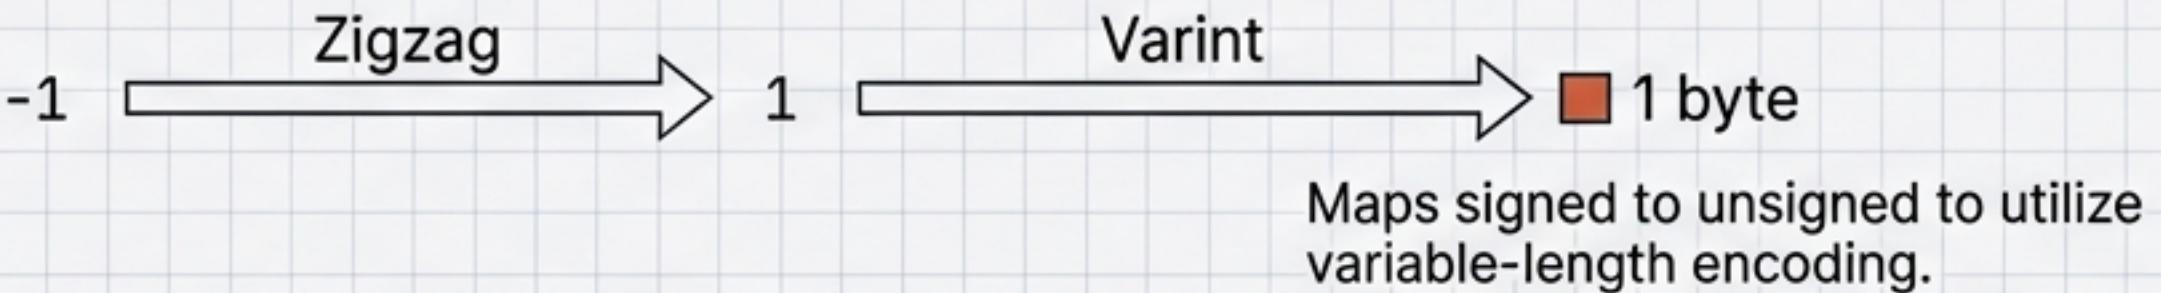


Floats
Gorilla XOR
8-12x Ratio

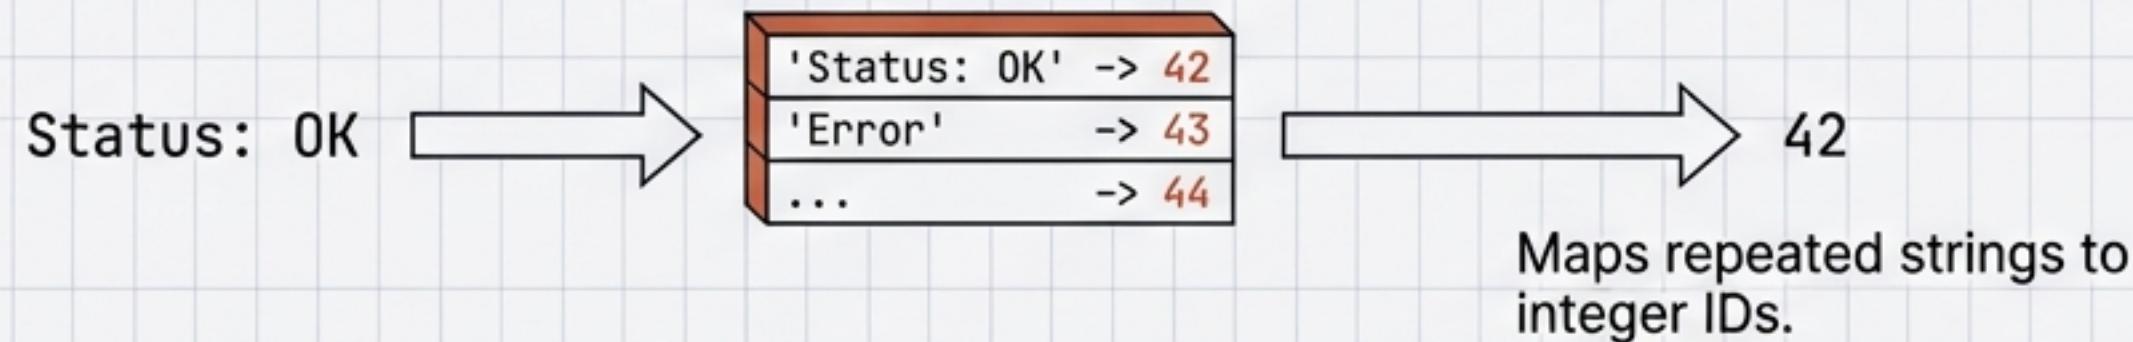
24.5 (01000001110001000...) \otimes (01000001110001000...) \rightarrow 1100010...

Facebook's Gorilla algorithm.
Strips leading/trailing zeros.

Integers
Zigzag + Varint
4-8x Ratio



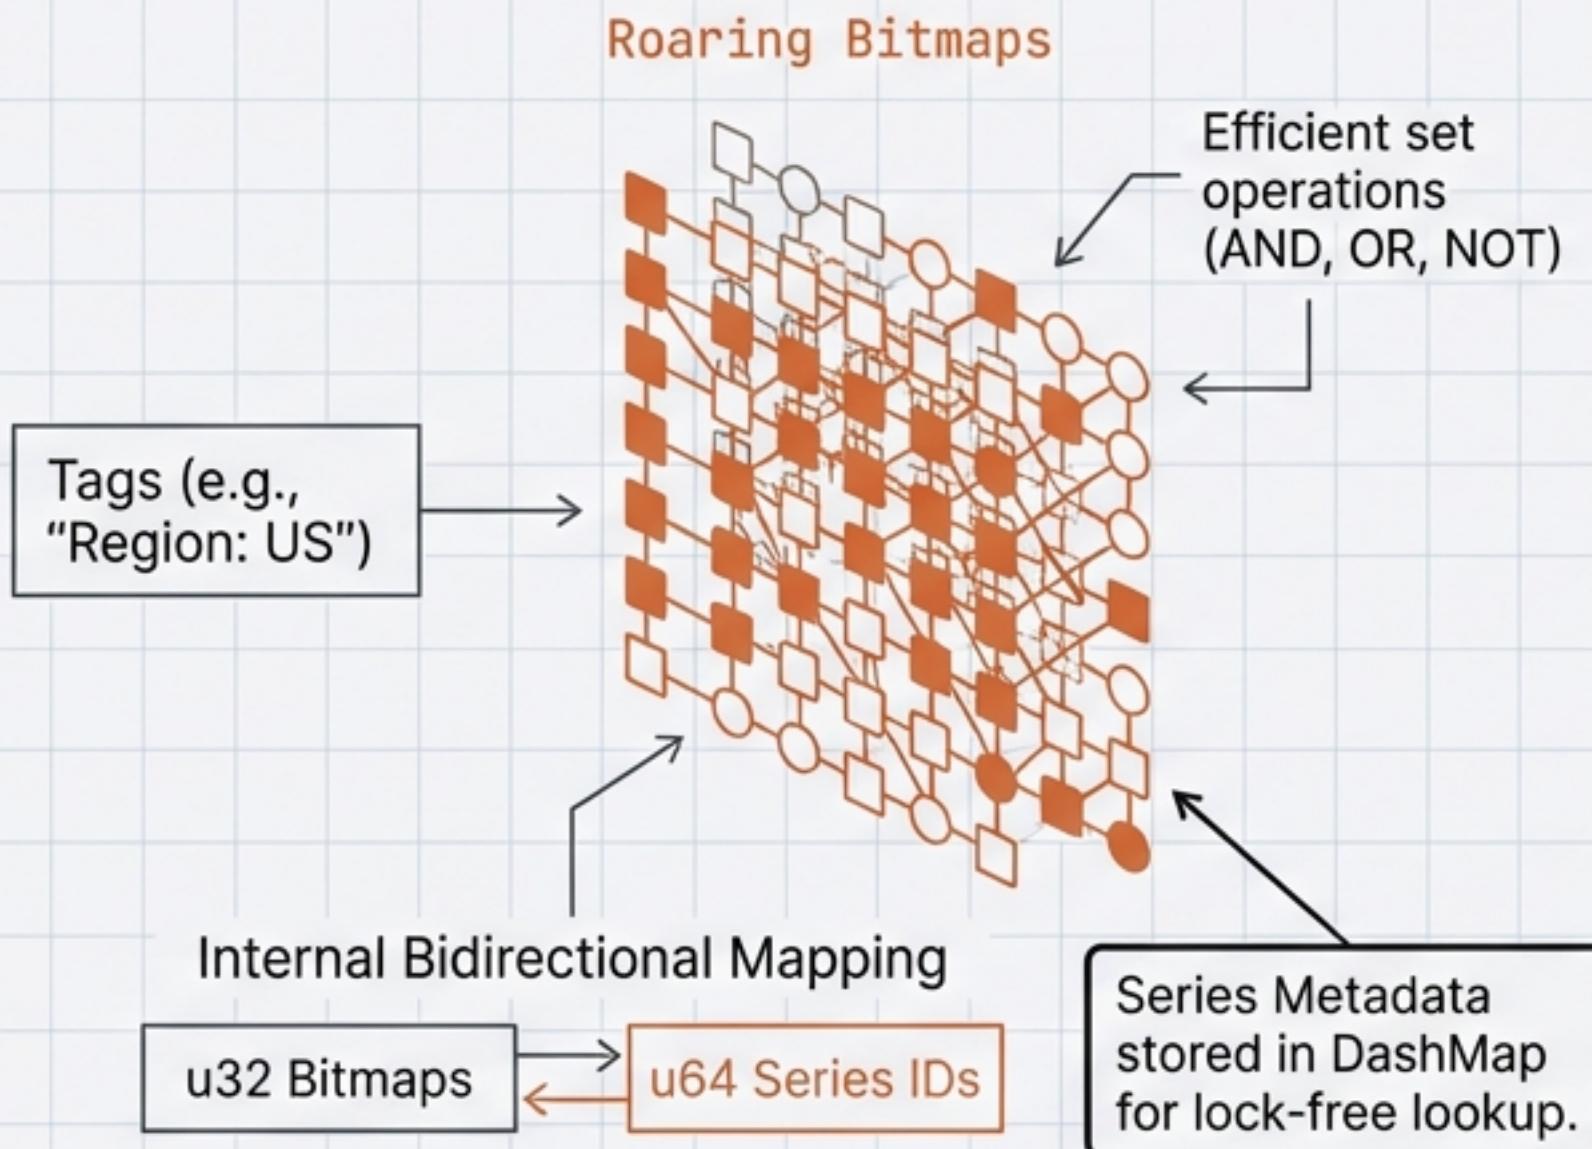
Strings
Dictionary Encoding
3-10x Ratio



High-Performance Indexing and Pruning

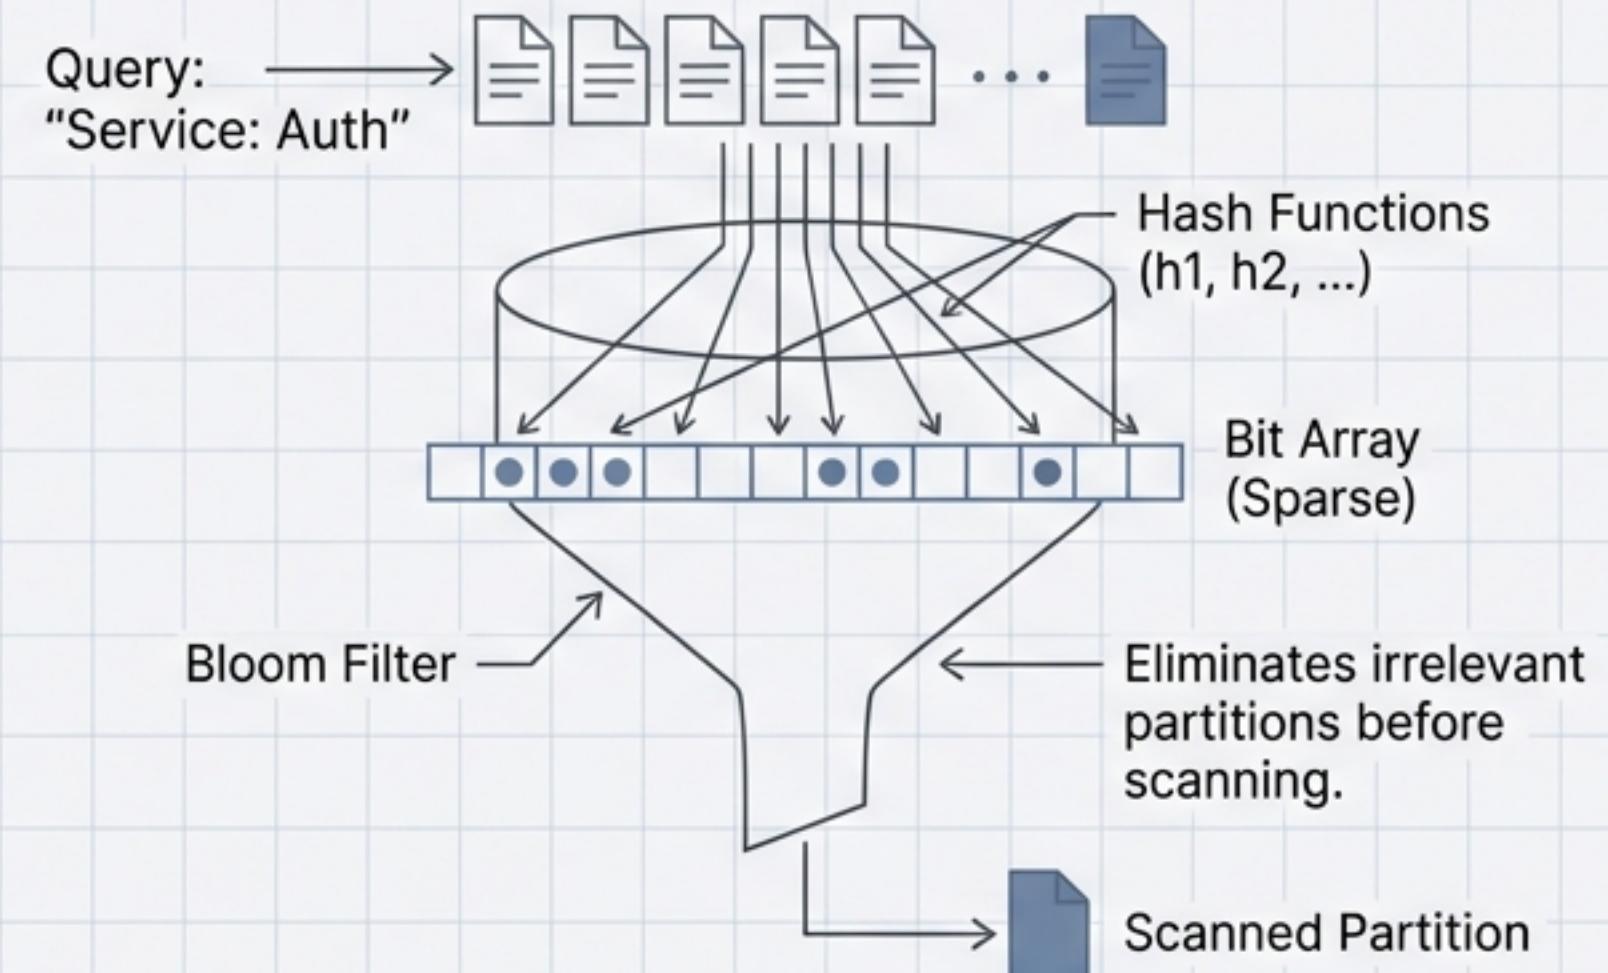
The Inverted Tag Index

Roaring Bitmaps



Partition Pruning

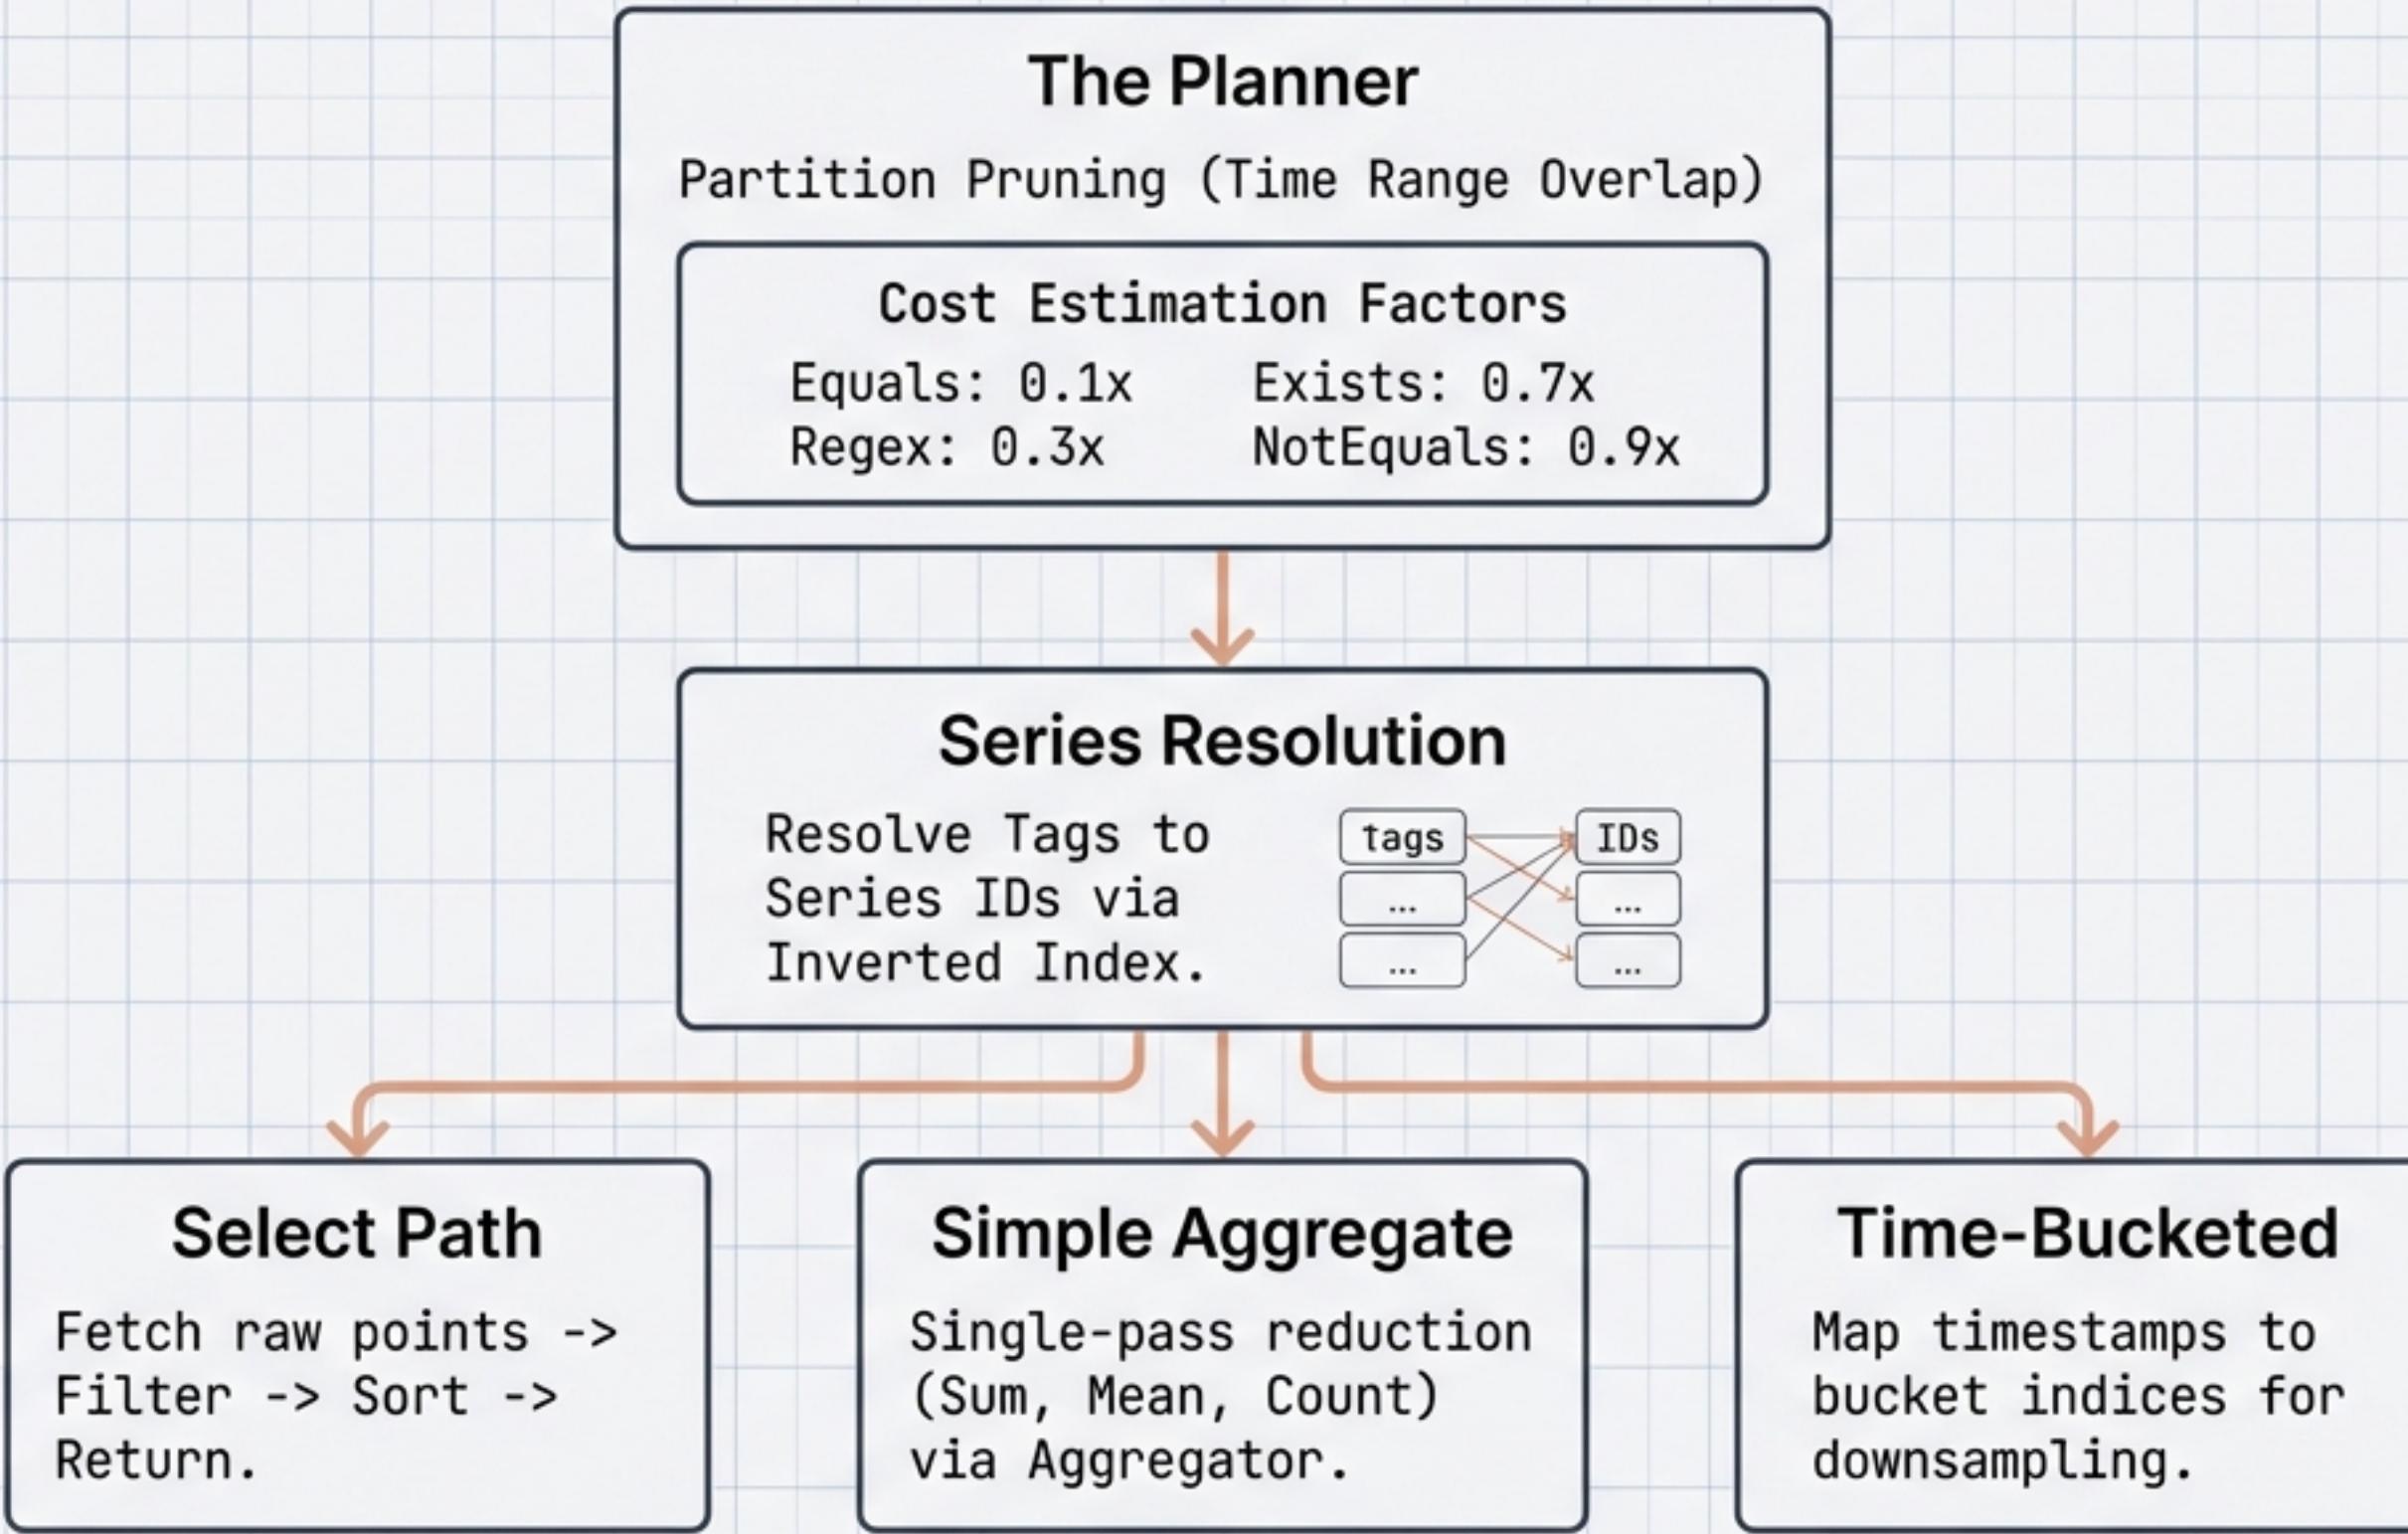
Bloom Filters



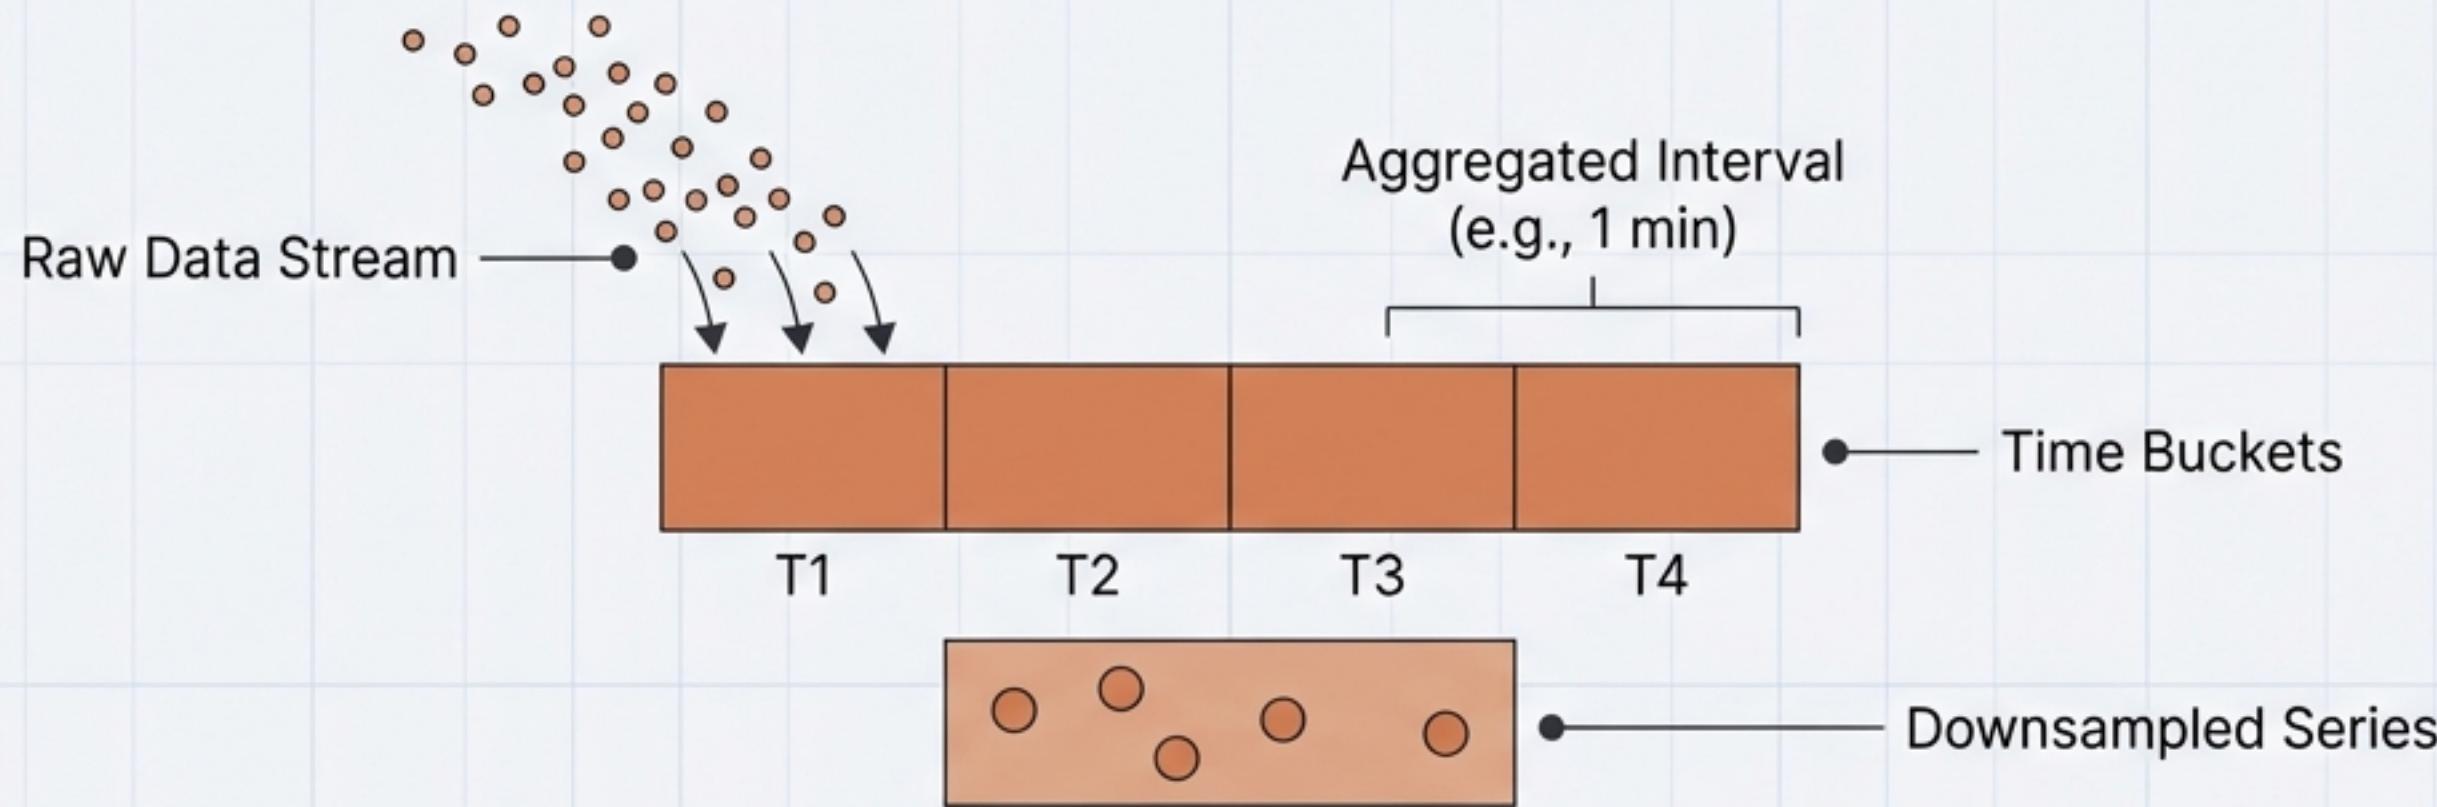
Specs

- **Hashing:** Double hashing $h(i) = h_1 + i * h_2$
- **Efficiency:** ~10 bits per item (1% False Positive Rate)

Query Planning and Execution Pipeline



Aggregation and Downsampling



Continuous Aggregates

Materialized view-like functionality. Decoupled from storage.

Logic:

Trigger: `current_time - last_processed > interval`

```
add_aggregation().group_by().retention()
```

Supported Functions

Standard: Count, Sum, Mean, Min, Max

Statistical: StdDev, Variance

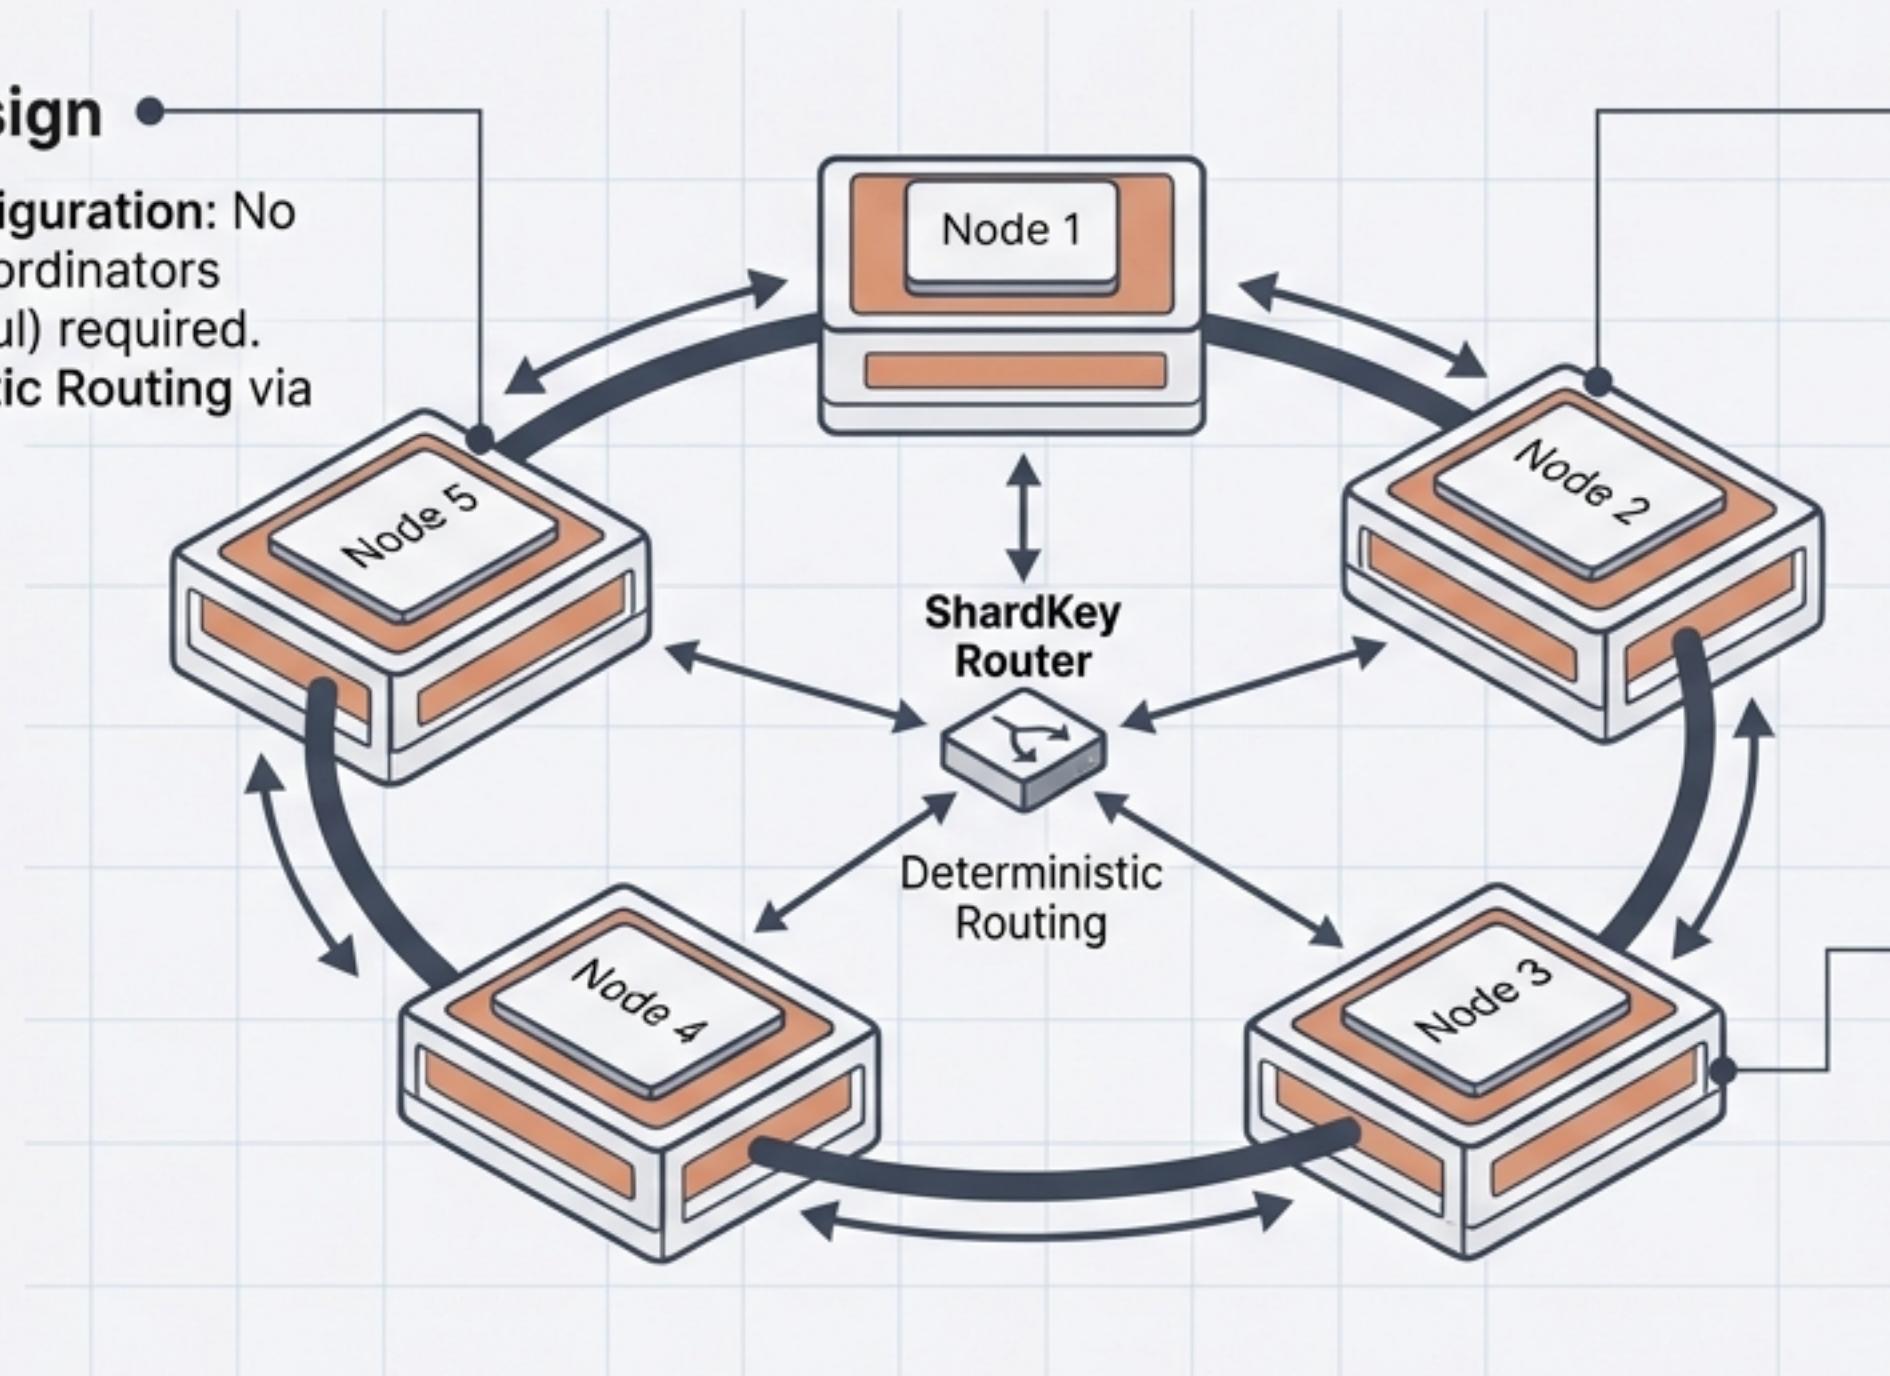
Advanced: Percentile (Linear Interpolation), First/Last

Thread-safe access via `parking_lot::RwLock`. Computation happens in `rusts-query`.

Clustering and Horizontal Scale

Cluster Design

- Static Configuration: No external coordinators (etcd/Consul) required.
- Deterministic Routing via ShardKey.



Sharding Strategies

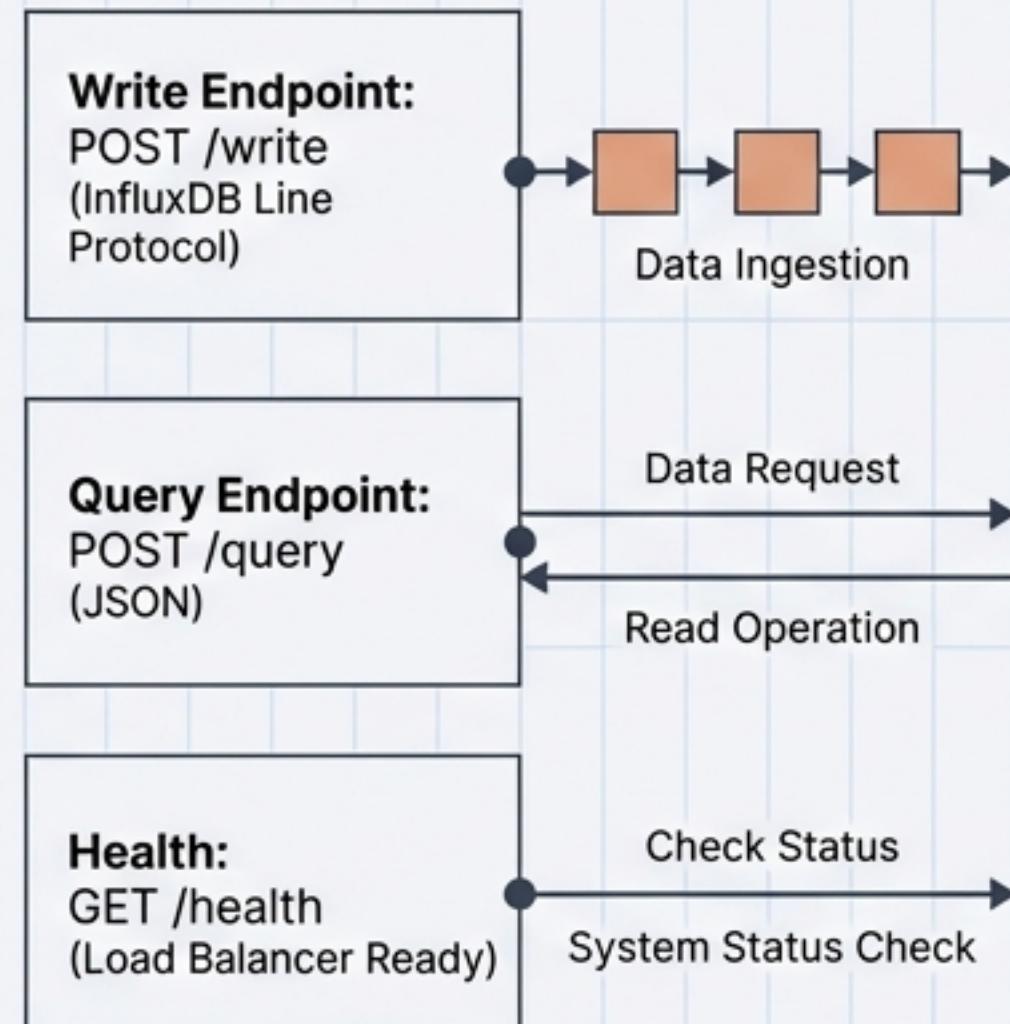
BySeries	Hash(measurement + tags) % shards. (Even distribution)
ByMeasurement	Isolation by measurement.
ByTime	Optimized for range queries.
Composite	SeriesID + TimeBucket. (Balanced)

Replication Modes

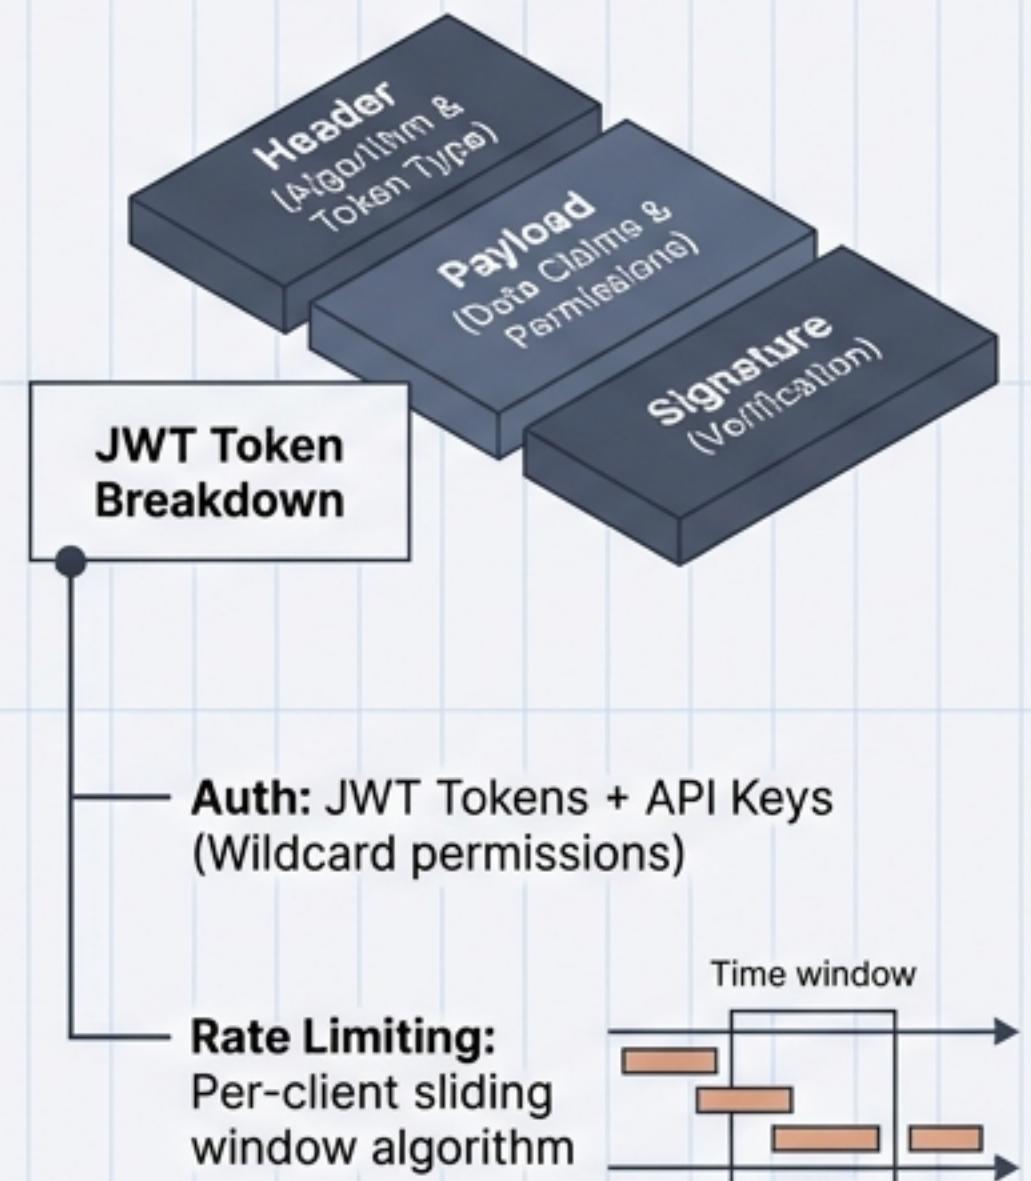
- **Sync:** Strong consistency. N acks. (High Latency)
- **Quorum:** Eventual consistency. $(N/2)+1$ acks.
- **Async:** Fire-and-forget. 0 acks. (Max Throughput)

The API Layer: Surface and Security

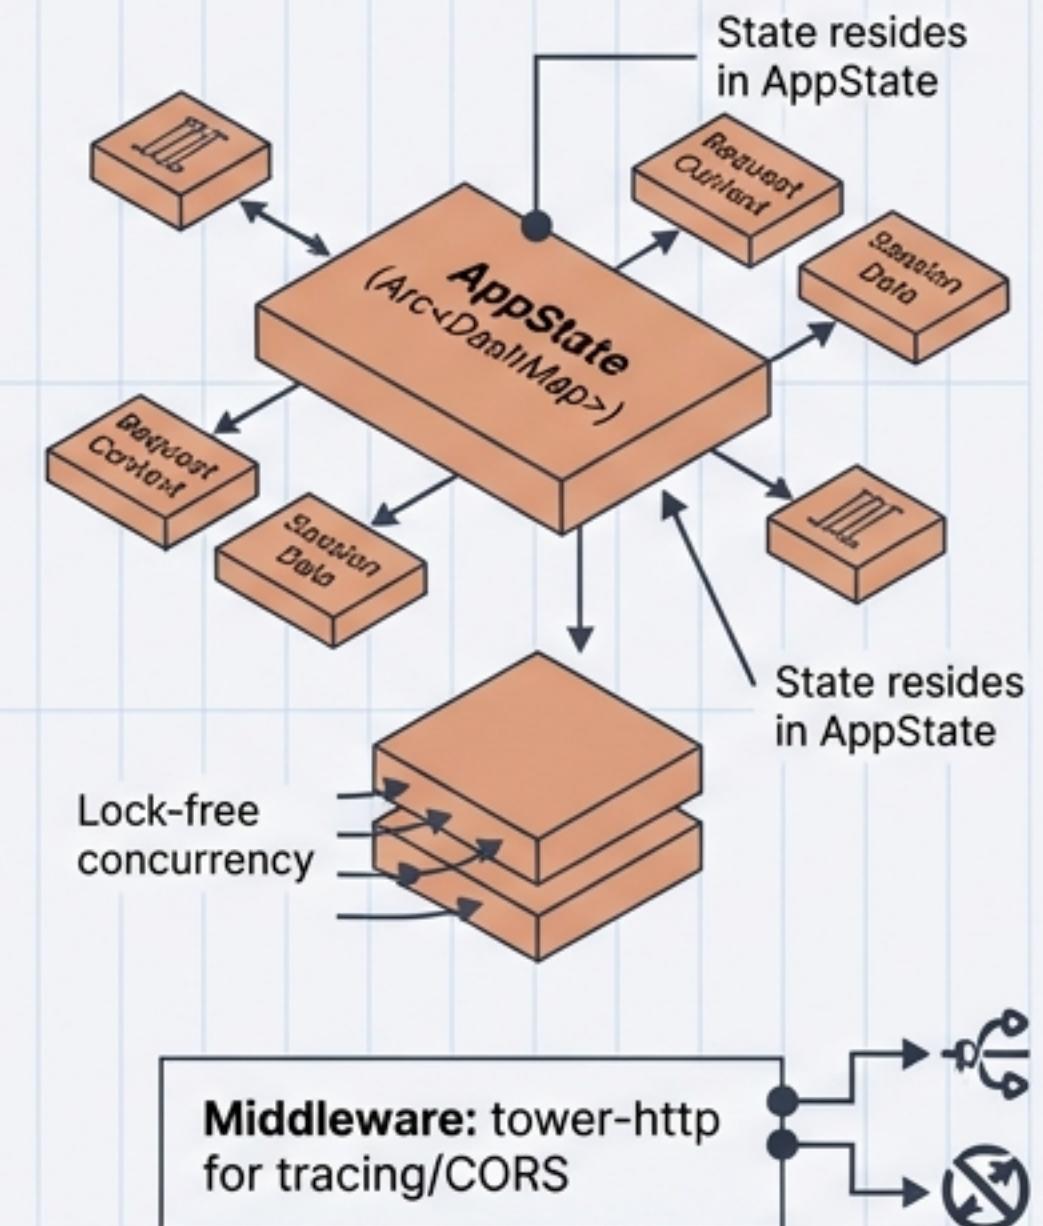
Interaction Model



Security & Auth



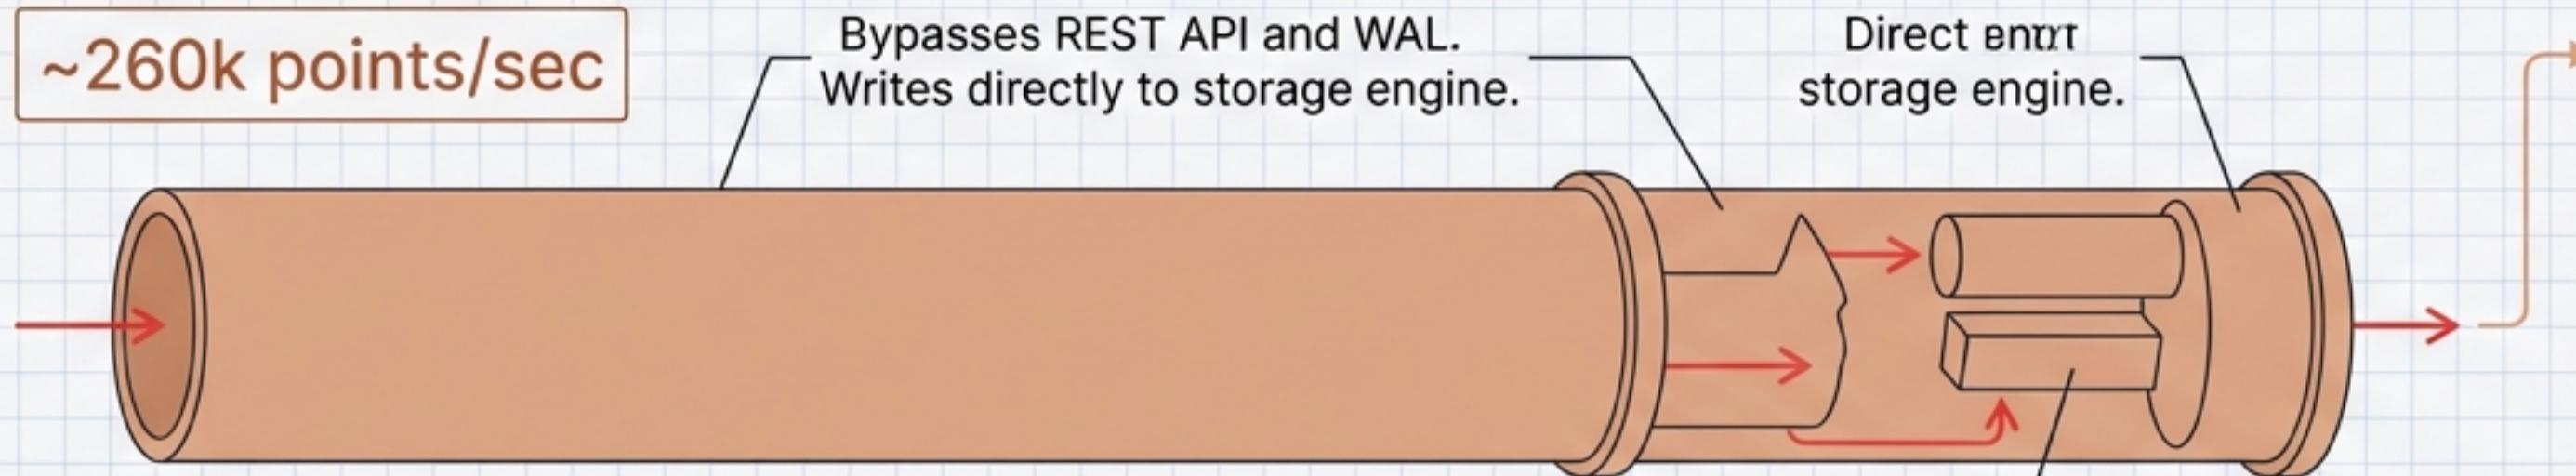
Stateless Design



High-Throughput Data Import

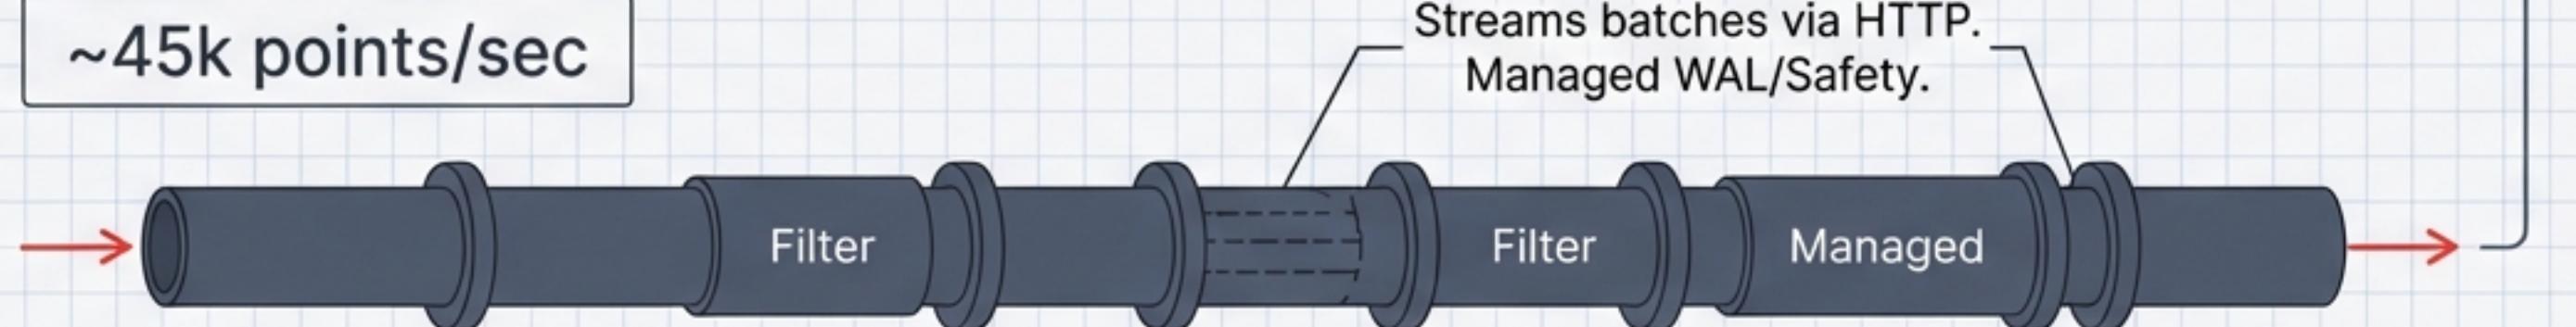
Direct Mode (The Performance King)

~260k points/sec



REST Mode (The Safe Option)

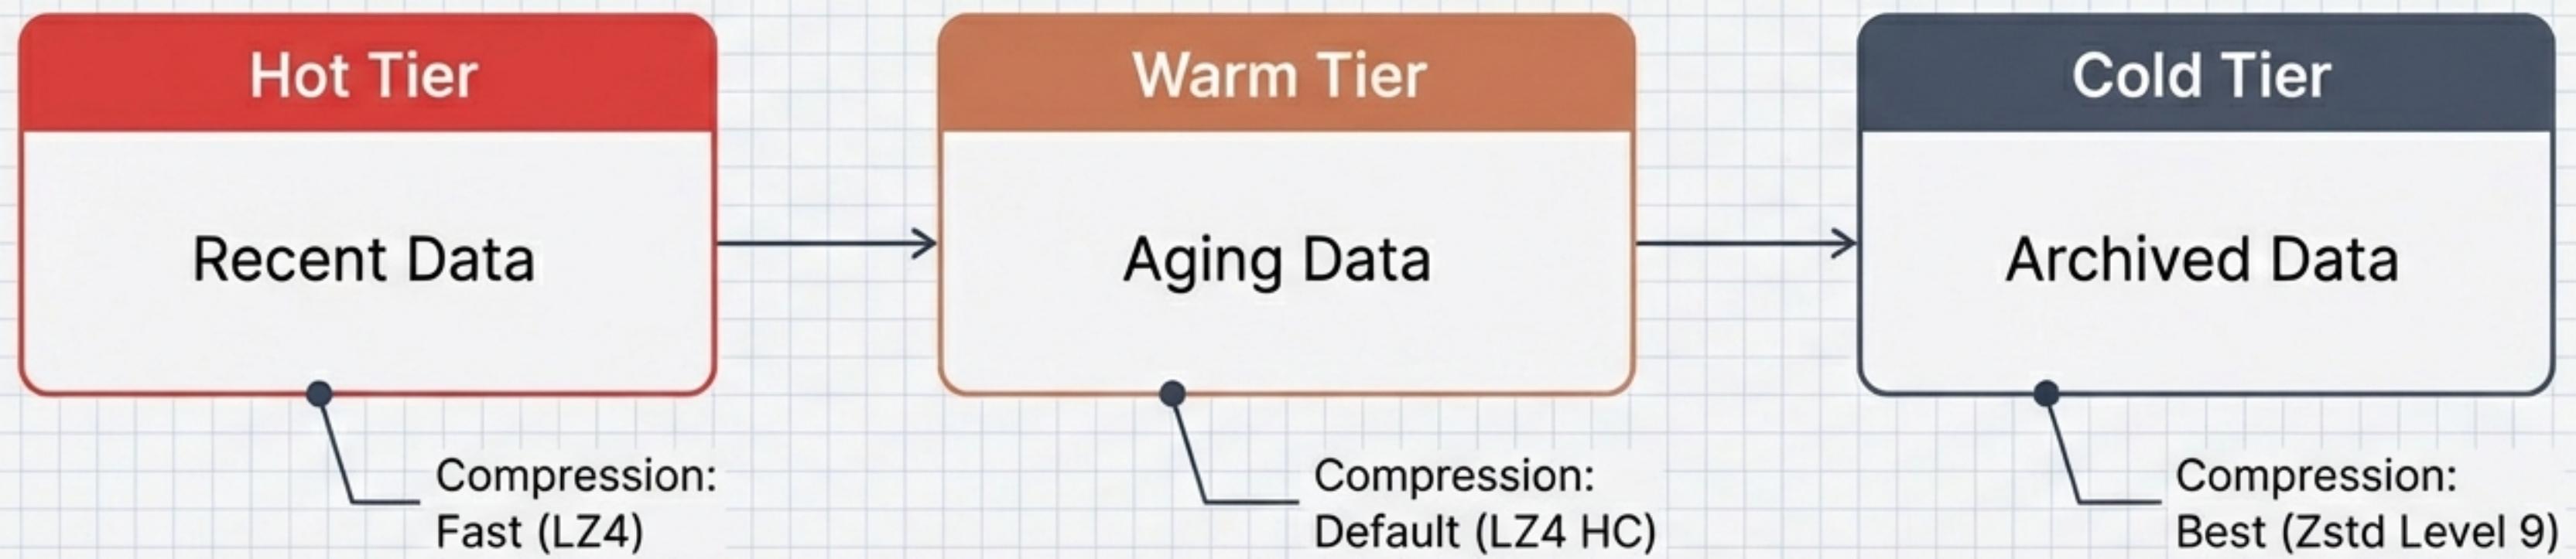
~45k points/sec



Deduplication

- Direct Mode:
Dedups by
Field only.
- REST Mode:
Dedups by
Field + Tag.

Data Lifecycle Management



Retention Policies

- Resolution Priority: Exact Match > Specific Pattern > Default.
- Pattern Matching: “cpu” (Exact), “cpu*” (Prefix), “*” (Wildcard).

Configuration and Tuning

rusts.yml

```
storage:  
  wal_durability: periodic # 100ms sync  
  partition_duration: 24h
```

Critical tradeoff knob.
"every_write" vs
"periodic".

```
memtable:  
  max_size: 64MB  
  max_points: 1000000  
  max_age: 60s
```

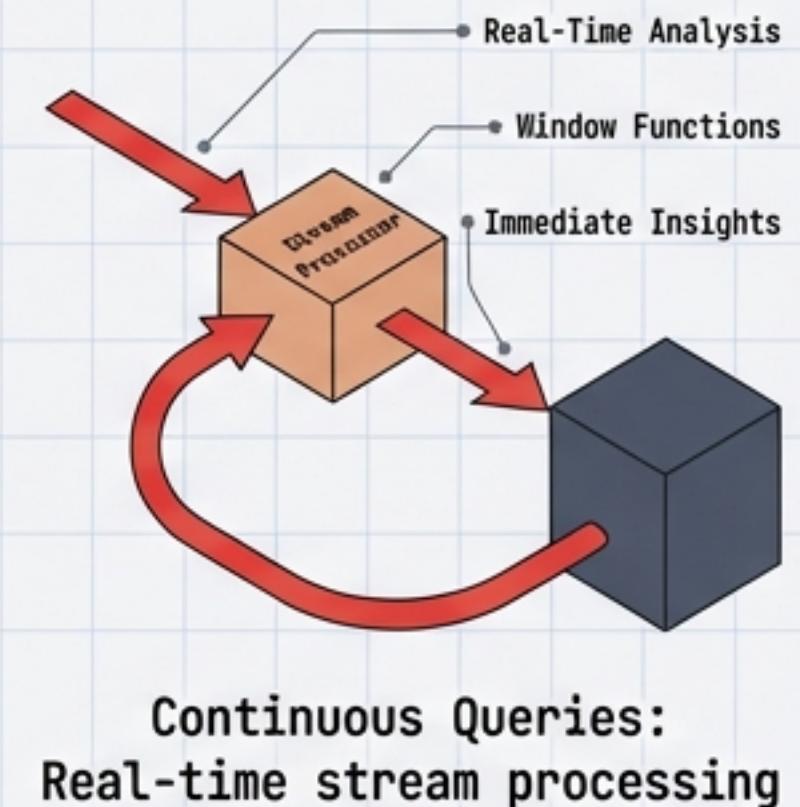
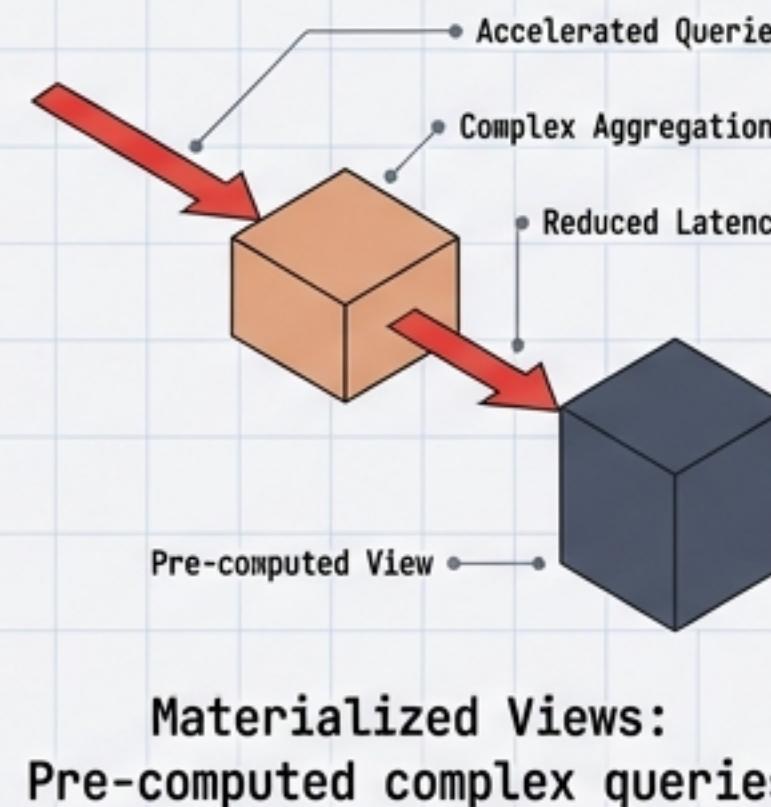
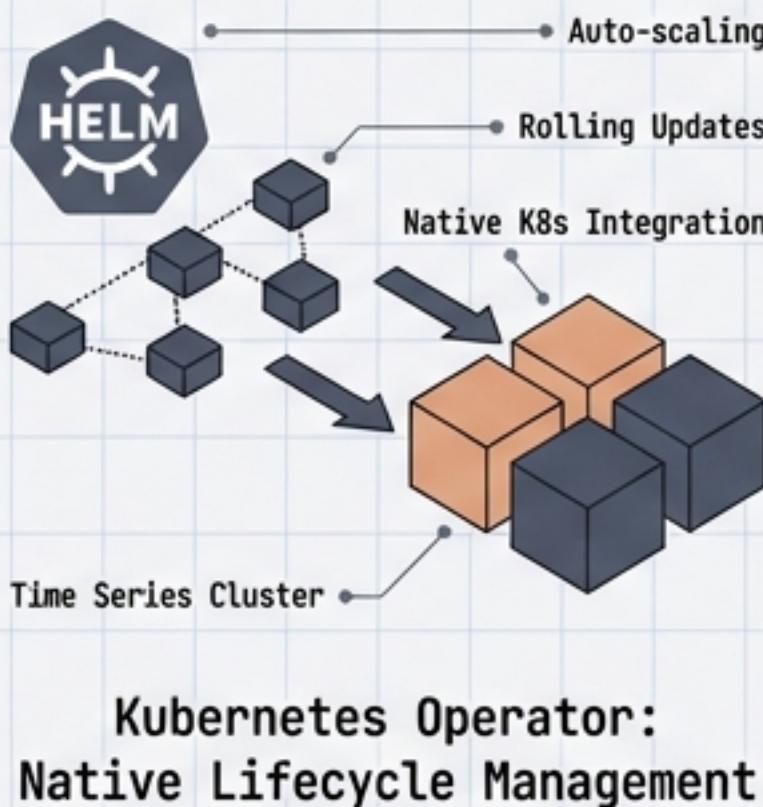
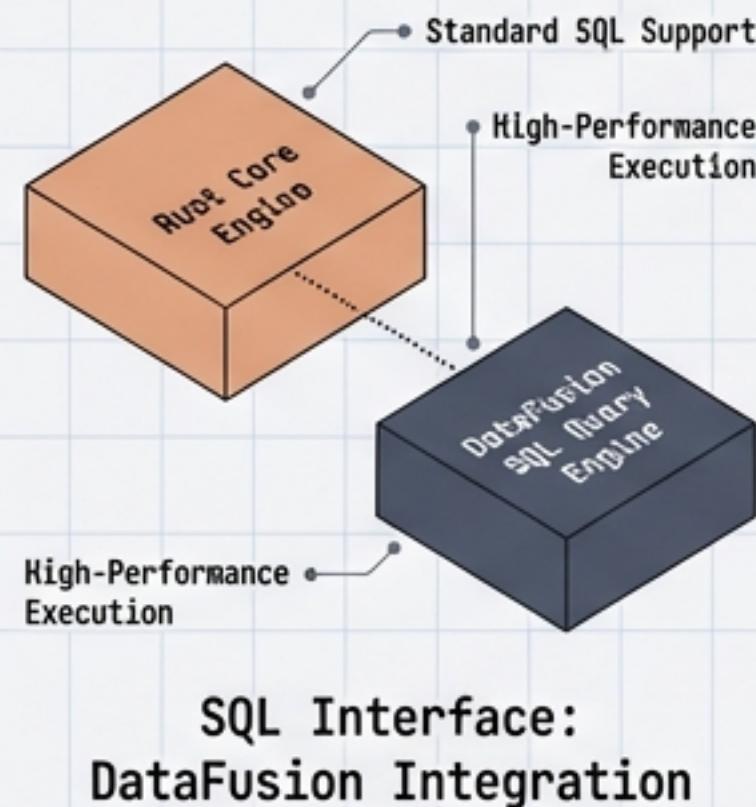
Trigger logic. Smaller
values reduce memory
pressure but increase
I/O fragmentation.

Adjust based on query patterns.

Default Network: Binds to 0.0.0.0:8086. Max body 10MB.

Roadmap and Future Horizons

Upcoming Features (Wireframe Visuals)



Getting Started

Prerequisites:
Rust 1.75+ and Cargo

> cargo run --release
Precision engineered for the future of time series.